


Lecture 10 cont'd: NP and Computational Intractability

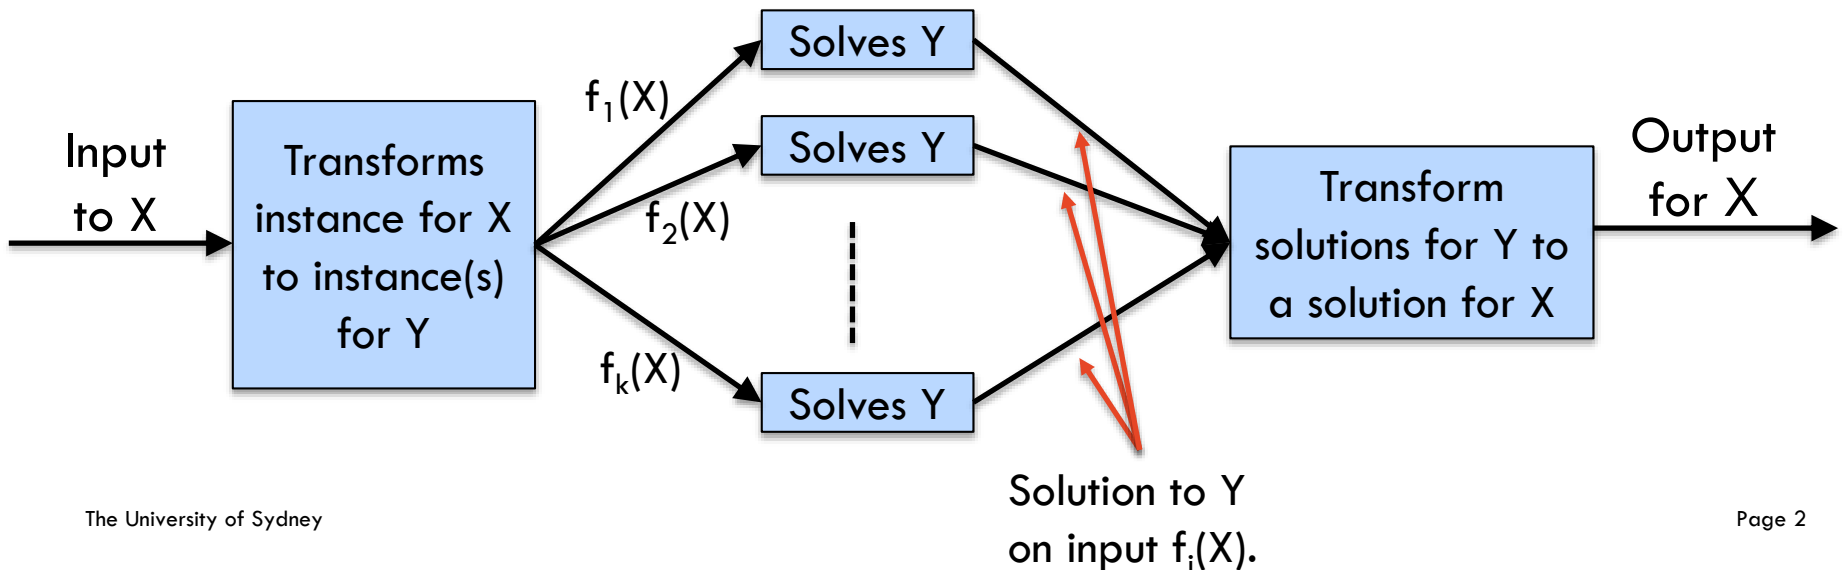


Polynomial-Time Reduction

Suppose we could solve problem Y in polynomial-time. What else could we solve in polynomial time?

Reduction. Problem X **polynomial reduces to** problem Y, denoted $X \leq_p Y$, if arbitrary instances of problem X can be solved using:

- Polynomial number of standard computational steps, plus
- Polynomial number of calls to an oracle that solves problem Y.



Polynomial-Time Reduction

Purpose. Classify problems according to **relative** difficulty.

1. **Design algorithms.** If $X \leq_p Y$ and Y can be solved in polynomial-time, then X **can** also be solved in polynomial time.
2. **Establish intractability.** If $X \leq_p Y$ and X cannot be solved in polynomial-time, then Y **cannot** be solved in polynomial time.

Summary – Lecture 10

- Polynomial time reductions

$3\text{-SAT} \leq_p \text{DIR HAMILTONIAN CYCLE} \leq_p \text{HAMILTONIAN CYCLE} \leq_p \text{TSP}$

$3\text{-SAT} \leq_p \text{INDEPENDENT-SET} \leq_p \text{VERTEX-COVER} \leq_p \text{SET-COVER}$

- Complexity classes:

P: Decision problems for which there is a **poly-time algorithm**.

NP: Decision problems for which there is a **poly-time certifier**.

NP-complete: A problem in NP such that every problem in NP polynomial reduces to it.

NP-hard: A problem such that every problem in NP polynomial reduces to it.

- Lots of problems are NP-complete

See <https://www.nada.kth.se/~viggo/wwwcompendium/>


Class NP-hard

Class NP-complete: A problem in NP such that every problem in NP polynomially reduces to it.


Class NP-hard:

A decision problem such that every problem in NP polynomially reduces to it.

not necessarily in NP



Many classes?



8.5 Sequencing Problems

Six basic genres


- Packing problems: SET-PACKING, INDEPENDENT SET.
- Covering problems: SET-COVER, VERTEX-COVER.
- Constraint satisfaction problems: SAT, 3-SAT.
- Sequencing problems: HAMILTONIAN-CYCLE, TSP.

$3\text{-SAT} \leq_p \text{DIR HAMILTONIAN CYCLE} \leq_p \text{HAMILTONIAN CYCLE} \leq_p \text{TSP}$

- Partitioning problems: 3D-MATCHING, 3-COLOR.
- Numerical problems: SUBSET-SUM, KNAPSACK.


Hamiltonian Cycle

HAM-CYCLE: given an undirected graph $G = (V, E)$, does there exist a simple cycle Γ that contains every node in V .



Hamiltonian Cycle


HAM-CYCLE: given an undirected graph $G = (V, E)$, does there exist a simple cycle Γ that contains every node in V .



HAM-CYCLE \in NP

Directed Hamiltonian Cycle

DIR-HAM-CYCLE: Given a directed graph $G = (V, E)$, does there exist a simple directed cycle Γ that contains every node in V ?




DIR-HAM-CYCLE \in NP

Directed Hamiltonian Cycle

DIR-HAM-CYCLE: Given a directed graph $G = (V, E)$, does there exist a simple directed cycle Γ that contains every node in V ?


Theorem: $\text{DIR-HAM-CYCLE} \leq_p \text{HAM-CYCLE}$.

Proof idea: Given a directed graph $G = (V, E)$, construct an undirected graph G' with $3n$ vertices.



Directed Hamiltonian Cycle


Claim: G has a Hamiltonian cycle iff G' does.



Proof:

- \Rightarrow – Suppose G has a directed Hamiltonian cycle Γ .
- Then G' has an undirected Hamiltonian cycle (same order).
- \Leftarrow – Suppose G' has an undirected Hamiltonian cycle Γ' .
- Γ' must visit nodes in G' using one of two orders:
 - ..., $B, G, R, B, G, R, B, G, R, B, \dots$
 - ..., $B, R, G, B, R, G, B, R, G, B, \dots$
 - Blue nodes in Γ' make up directed Hamiltonian cycle Γ in G , or reverse of one. ▀

$\text{DIR-HAM-CYCLE} \leq_p \text{HAM-CYCLE}$



3-SAT Reduces to Directed Hamiltonian Cycle

Theorem: $3\text{-SAT} \leq_p \text{DIR-HAM-CYCLE}$.

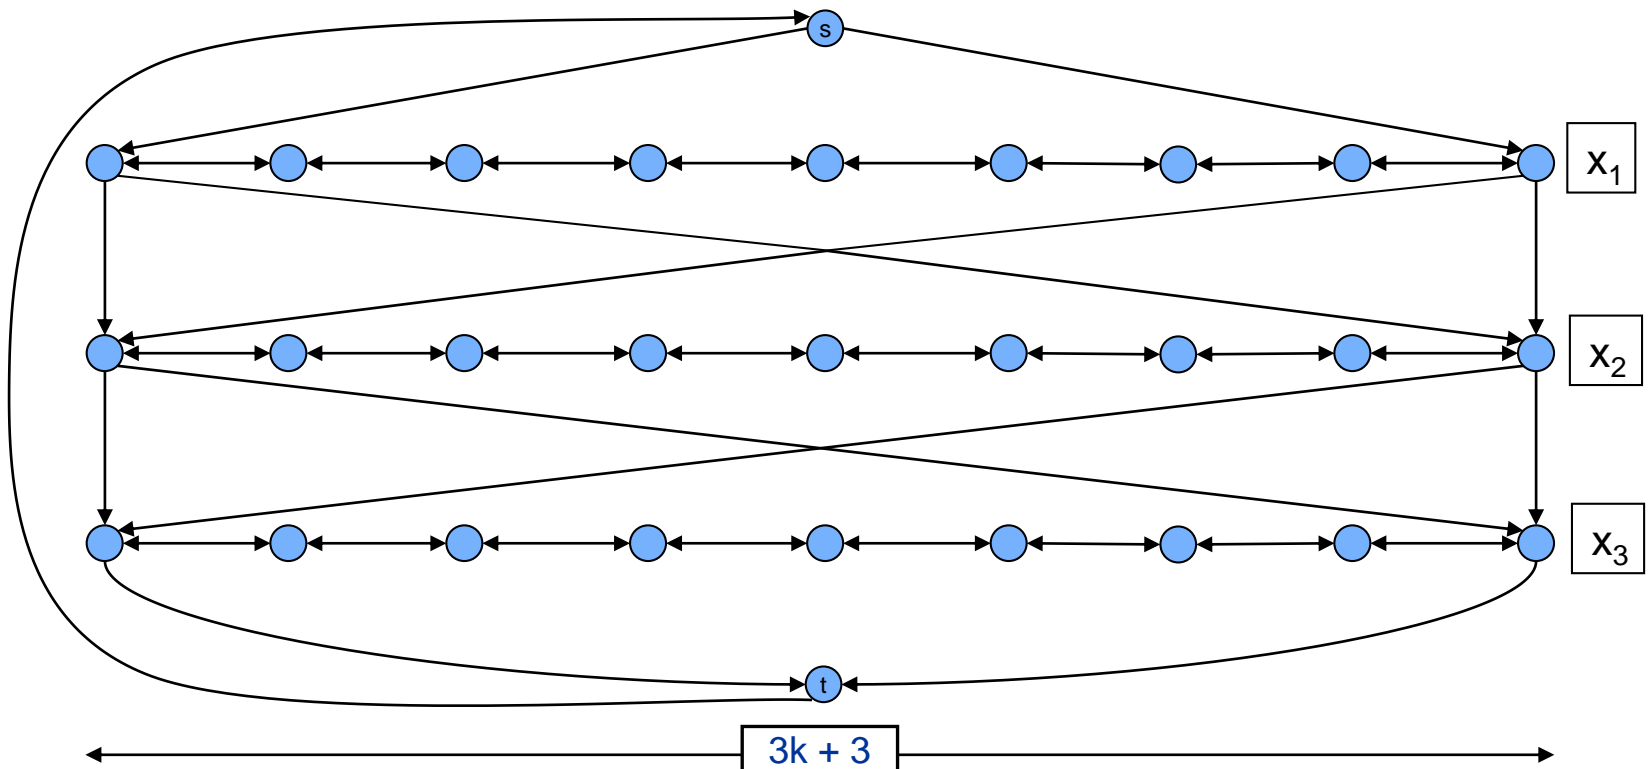
Proof: Given an instance Φ of 3-SAT, we construct an instance of DIR-HAM-CYCLE that has a Hamiltonian cycle iff Φ is satisfiable.

Construction. First, create graph that has 2^n Hamiltonian cycles which correspond in a natural way to 2^n possible truth assignments.

3-SAT Reduces to Directed Hamiltonian Cycle


Construction: Given a 3-SAT instance Φ with n variables x_i and k clauses.

- Construct G to have 2^n Hamiltonian cycles.
- **Intuition:** Traverse path i from left to right \Leftrightarrow set variable $x_i = 1$.




3-SAT Reduces to Directed Hamiltonian Cycle

- Construction. Given 3-SAT instance Φ with n variables x_i and k clauses.
 - For each clause: add a node and 6 edges.




3-SAT Reduces to Directed Hamiltonian Cycle

- Construction. Given 3-SAT instance Φ with n variables x_i and k clauses.
 - For each clause: add a node and 6 edges.




3-SAT Reduces to Directed Hamiltonian Cycle

- Construction. Given 3-SAT instance Φ with n variables x_i and k clauses.
 - For each clause: add a node and 6 edges.



3-SAT Reduces to Directed Hamiltonian Cycle

- Construction. Given 3-SAT instance Φ with n variables x_i and k clauses.
 - For each clause: add a node and 6 edges.




3-SAT Reduces to Directed Hamiltonian Cycle

Claim: Φ is satisfiable iff G has a Hamiltonian cycle.

Proof: \Rightarrow

- Suppose 3-SAT instance has satisfying assignment x^* .
- Then, define Hamiltonian cycle in G as follows:
 - if $x_i^* = 1$, traverse row i from left to right
 - if $x_i^* = 0$, traverse row i from right to left
 - for each clause C_i , there will be at least one row i in which we are going in "correct" direction to splice node C_i into tour



3-SAT Reduces to Directed Hamiltonian Cycle

Claim: Φ is satisfiable iff G has a Hamiltonian cycle.

Proof: \Leftarrow

- Suppose G has a Hamiltonian cycle Γ .
- If Γ enters clause node C_i , it must depart on mate edge.
 - thus, nodes immediately before and after C_i are connected by an edge e in G
 - removing C_i from cycle, and replacing it with edge e yields Hamiltonian cycle on $G \setminus \{C_i\}$
- Continuing in this way, we are left with Hamiltonian cycle Γ' in $G - \{C_1, C_2, \dots, C_k\}$.
- Set $x_i^* = 1$ iff Γ' traverses row i left to right.
- Since Γ visits each clause node C_i , at least one of the paths is traversed in "correct" direction, and each clause is satisfied. ▀

3-SAT Reduces to Directed Hamiltonian Cycle

Theorem: $3\text{-SAT} \leq_p \text{DIR-HAM-CYCLE}$.

- 3-SAT is NP-complete
- $\text{DIR-HAM-CYCLE} \in \text{NP}$

\Rightarrow DIR-HAM-CYCLE is NP-complete

Directed HAM-CYCLE reduces to HAM-CYCLE


Theorem: $\text{DIR-HAM-CYCLE} \leq_p \text{HAM-CYCLE}$.

- DIR-HAM-CYCLE is NP-complete
- HAM-CYCLE \in NP

\Rightarrow HAM-CYCLE is NP-complete

Travelling Salesperson Problem

TSP: Given a set of n cities and a pairwise distance function $d(u, v)$, is there a tour of length $\leq D$?




All 13,509 cities in US with a population of at least 500

Reference: <http://www.tsp.gatech.edu>

Travelling Salesperson Problem

TSP: Given a set of n cities and a pairwise distance function $d(u, v)$, is there a tour of length $\leq D$?



Optimal TSP tour
Reference: <http://www.tsp.gatech.edu>

Travelling Salesperson Problem

TSP: Given a set of n cities and a pairwise distance function $d(u, v)$, is there a tour of length $\leq D$?




11,849 holes to drill in a programmed logic array

Reference: <http://www.tsp.gatech.edu>

Travelling Salesperson Problem

TSP: Given a set of n cities and a pairwise distance function $d(u, v)$, is there a tour of length $\leq D$?



Optimal TSP tour
Reference: <http://www.tsp.gatech.edu>

Travelling Salesperson Problem

TSP: Given a set of n cities and a pairwise distance function $d(u, v)$, is there a tour of length $\leq D$?

HAM-CYCLE: given a graph $G = (V, E)$, does there exist a simple cycle that contains every node in V ?

Theorem: $\text{HAM-CYCLE} \leq_p \text{TSP}$.

Proof:

- Given instance $G = (V, E)$ of HAM-CYCLE, create n cities with distance function

$$d(u, v) = \begin{cases} 1 & \text{if } (u, v) \in E \\ 2 & \text{if } (u, v) \notin E \end{cases}$$

- TSP instance has tour of length $\leq n$ iff G is Hamiltonian. ▀

TSP is NP-complete

Theorem: $\text{HAM-CYCLE} \leq_p \text{TSP}$.

- HAM-CYCLE is NP-complete
- $\text{TSP} \in \text{NP}$

\Rightarrow TSP (decision version) is NP-complete

NP-complete games and puzzles

- Battleship
- Candy Crush Saga
- Donkey Kong
- Eternity II
- FreeCell
- Lemmings
- Minesweeper Consistency Problem
- Pokémon
- SameGame
- Sudoku
- Tetris
- Rush Hour
- Hex
- Super Mario Bros

Summary

- Polynomial time reductions

$3\text{-SAT} \leq_p \text{DIR HAMILTONIAN CYCLE} \leq_p \text{HAMILTONIAN CYCLE} \leq_p \text{TSP}$

$3\text{-SAT} \leq_p \text{INDEPENDENT-SET} \leq_p \text{VERTEX-COVER} \leq_p \text{SET-COVER}$

- Complexity classes:

P: Decision problems for which there is a **poly-time algorithm**.

NP: Decision problems for which there is a **poly-time certifier**.

NP-complete: A problem in NP such that every problem in NP polynomial reduces to it.

NP-hard: A problem such that every problem in NP polynomial reduces to it.

- Lots of problems are NP-complete

See <https://www.nada.kth.se/~viggo/wwwcompendium/>

Lecture 11: Coping with hardness



THE UNIVERSITY OF
SYDNEY

Algorithms and hardness

Algorithmic techniques:

- Greedy algorithms [Lecture 3]
- Divide & Conquer algorithms [Lecture 4]
- Sweepline algorithms [Lecture 5]
- Dynamic programming algorithms [Lectures 6 and 7]
- Network flow algorithms [Lectures 8 and 9]

Hardness:

- NP-hardness [Lecture 10]: $O(n^c)$ algorithm is unlikely

Today

- How can we cope with hard problems?

Algorithms and hardness

Lots of problems that we can solve efficiently:

- MST
- Shortest path
- Scheduling
- Max flow
- ...

But lots of problems that we can't solve efficiently:

- All NP-complete problems: $O(n^c)$ algorithm is unlikely
vertex cover, independent set, 3-SAT,...
- But what if we need to solve an NP-complete problem?


Cope with NP-complete problems?



"I can't find an efficient algorithm, but neither can all these famous people."

Cope with NP-complete problems?

- Heuristics: Local search, simulated annealing, neural networks...
- **Randomized algorithms:** Not always correct, but a probability is guaranteed.
- **Approximation algorithms:** Not optimal solution, but within a guaranteed error.
- Fixed-parameter algorithms: Algorithms whose complexity depends on other parameters than n .
- Efficient exact algorithms: Euclidean TSP
- **Restricted instance:** trees, bipartite graphs...



Coping With NP-Completeness

Question: What should I do if I need to solve an NP-complete problem?

Answer: Theory says you're unlikely to find poly-time algorithm.

Must sacrifice one of three desired features.

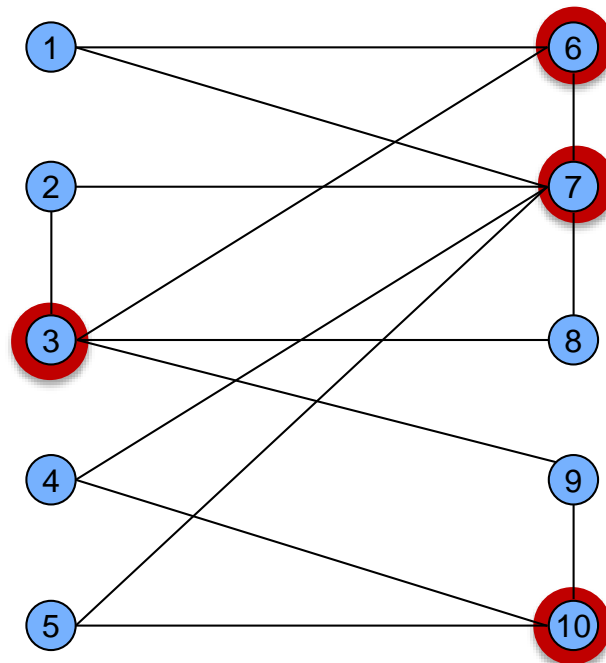
- Solve problem to optimality.
 - Approximation algorithms
 - Randomized algorithms
- Solve problem in polynomial time.
 - Exact exponential time algorithms
- Solve arbitrary instances of the problem.
 - Solve restricted classes of instances
 - Parametrized algorithms

10.1 Solving restricted instances

For example in the case when the vertex cover is small.

10.1 Finding Small Vertex Covers

VERTEX COVER: Given a graph $G = (V, E)$ and an integer k , is there a subset of vertices $S \subseteq V$ such that $|S| \leq k$, and for each edge (u, v) either $u \in S$, or $v \in S$, or both.



$$k = 4$$

$$S = \{ 3, 6, 7, 10 \}$$

Vertex Cover

VERTEX COVER: Given a graph $G = (V, E)$ and an integer k , is there a subset of vertices $S \subseteq V$ such that $|S| \leq k$, and for each edge (u, v) either $u \in S$, or $v \in S$, or both.

Vertex Cover (or Independent Set) arises naturally in many applications:

- dynamic detection of race conditions (distributed systems).
- computational biology
- Biochemistry
- Pattern recognition
- Computer vision
- ...

What should we do if we need to solve it?

Finding Small Vertex Covers

Question: What if k is small?

Brute force: $O(kn^{k+1})$.

- Try all $\binom{n}{k} \in O(n^k)$ subsets of size k .
- Takes $O(kn)$ time to check whether a subset is a vertex cover.

Finding Small Vertex Covers

Question: What if k is small?

Brute force: $O(kn^{k+1})$.

- Try all $\binom{n}{k} \in O(n^k)$ subsets of size k .
- Takes $O(kn)$ time to check whether a subset is a vertex cover.

Aim: Limit exponential dependency on k , e.g., to $O(2^k kn)$.

Example: $n = 1000, k = 10$.


- Brute force. $kn^{k+1} = 10^{34} \Rightarrow$ infeasible.
- Better. $2^k kn = 10^7 \Rightarrow$ feasible.

Remark. If k is a constant, algorithm is poly-time; if k is a small constant, then it's also practical.

Finding Small Vertex Covers

Theorem: Let (u,v) be an edge of G . G has a vertex cover of size $\leq k$ iff at least one of $G \setminus \{u\}$ and $G \setminus \{v\}$ has a vertex cover of size $\leq k-1$.

delete u and all incident edges



Finding Small Vertex Covers


Theorem: Let (u,v) be an edge of G . G has a vertex cover of size $\leq k$ iff at least one of $G \setminus \{u\}$ and $G \setminus \{v\}$ has a vertex cover of size $\leq k-1$.

delete u and all incident edges

Proof:

\Rightarrow

- Suppose G has a vertex cover S of size $\leq k$.
- S contains either u or v (or both). Assume it contains u .
- $S \setminus \{u\}$ is a vertex cover of $G \setminus \{u\}$.



Finding Small Vertex Covers

Theorem: Let (u,v) be an edge of G . G has a vertex cover of size $\leq k$ iff at least one of $G \setminus \{u\}$ and $G \setminus \{v\}$ has a vertex cover of size $\leq k-1$.

delete u and all incident edges


Proof:

\Rightarrow

- Suppose G has a vertex cover S of size $\leq k$.
- S contains either u or v (or both). Assume it contains u .
- $S \setminus \{u\}$ is a vertex cover of $G \setminus \{u\}$.

\Leftarrow

- Suppose S is a vertex cover of $G \setminus \{u\}$ of size $\leq k-1$.
- Then $S \cup \{u\}$ is a vertex cover of G of size k . ▀



Finding Small Vertex Covers

Theorem: Let (u,v) be an edge of G . G has a vertex cover of size $\leq k$ iff at least one of $G \setminus \{u\}$ and $G \setminus \{v\}$ has a vertex cover of size $\leq k-1$.

delete u and all incident edges


Proof:

\Rightarrow

- Suppose G has a vertex cover S of size $\leq k$.
- S contains either u or v (or both). Assume it contains u .
- $S \setminus \{u\}$ is a vertex cover of $G \setminus \{u\}$.

\Leftarrow

- Suppose S is a vertex cover of $G \setminus \{u\}$ of size $\leq k-1$.
- Then $S \cup \{u\}$ is a vertex cover of G of size k . ▀





Observation: If G has a vertex cover of size k , it has $\leq k(n-1)$ edges.

Proof: Each vertex covers at most $n-1$ edges. ▀

Finding Small Vertex Covers: Algorithm

Theorem: The following algorithm determines if G has a vertex cover of size $\leq k$ in $O(2^k kn)$ time.



```
boolean Vertex-Cover( $G, k$ ) {  
    if ( $G$  contains no edges)    return true  
    if ( $G$  contains  $> k(n-1)$  edges) return false  
  
    let  $(u, v)$  be any edge of  $G$   
     $a = \text{Vertex-Cover}(G \setminus \{u\}, k-1)$   
     $b = \text{Vertex-Cover}(G \setminus \{v\}, k-1)$   
    return  $a$  or  $b$   
}
```



Finding Small Vertex Covers: Algorithm

Theorem: The following algorithm determines if G has a vertex cover of size $\leq k$ in $O(2^k kn)$ time.



```
boolean Vertex-Cover( $G, k$ ) {  
    if ( $G$  contains no edges)    return true  
    if ( $G$  contains  $> k(n-1)$  edges) return false  
  
    let  $(u, v)$  be any edge of  $G$   
     $a = \text{Vertex-Cover}(G \setminus \{u\}, k-1)$   
     $b = \text{Vertex-Cover}(G \setminus \{v\}, k-1)$   
    return  $a$  or  $b$   
}
```



Finding Small Vertex Covers: Algorithm

Theorem: The following algorithm determines if G has a vertex cover of size $\leq k$ in $O(2^k kn)$ time.

```
boolean Vertex-Cover( $G, k$ ) {  
    if ( $G$  contains no edges)    return true  
    if ( $G$  contains  $> k(n-1)$  edges) return false  
  
    let  $(u, v)$  be any edge of  $G$   
     $a = \text{Vertex-Cover}(G \setminus \{u\}, k-1)$   
     $b = \text{Vertex-Cover}(G \setminus \{v\}, k-1)$   
    return  $a$  or  $b$   
}
```



Proof:



- Correctness follows from previous two claims.
- There are $\leq 2^{k+1}$ nodes in the recursion tree; each invocation takes $O(kn)$ time. ▀

Finding Small Vertex Covers: Algorithm

Theorem: Vertex cover can be solved in $O(2^k kn)$ time.

This is fine as long as k is (a small) constant.

What if k is not a small constant?



10.2 Solving NP-Hard Problems on restricted input instances

For example special cases of graphs:

- trees,
- bipartite graphs,
- planar graphs,
- ...

Independent Set on Trees


INDEPENDENT-SET: Given a graph $G = (V, E)$ and an integer k , is there a subset of vertices $S \subseteq V$ such that $|S| \geq k$, and for each edge at most one of its endpoints is in S ?

Problem: Given a **tree**, find a maximum IS.

Key observation: If v is a leaf, there exists a maximum size independent set containing v .

Proof: [exchange argument]


- Consider a max cardinality independent set S .
- If $v \in S$, we're done.
- If $u \notin S$ and $v \notin S$, then $S \cup \{v\}$ is independent $\Rightarrow S$ not maximum.
- If $u \in S$ and $v \notin S$, then $S \cup \{v\}/\{u\}$ is independent. (exchange) ▪



Independent Set on Trees: Greedy Algorithm

Theorem: The following greedy algorithm finds a maximum cardinality independent set in forests (and hence trees).

```
Independent-Set-In-A-Forest(F) {  
  S  $\leftarrow$   $\emptyset$   
  while (F has at least one edge) {  
    Let  $e = (u,v)$  be an edge such that  $v$  is a leaf  
    Add  $v$  to S  
    Delete from F nodes  $u$  and  $v$ , and all edges  
      incident to them.  
  }  
  return S  
}
```



Proof: Correctness follows from the previous key observation. ▀

Remark. Can implement in $O(n)$ time by considering nodes in postorder.


Independent Set on Trees

INDEPENDENT-SET: Given a graph $G = (V, E)$ and an integer k , is there a subset of vertices $S \subseteq V$ such that $|S| \geq k$, and for each edge at most one of its endpoints is in S ?

Problem: Given a **tree**, find a maximum IS.


Theorem:

INDEPENDENT-SET on trees can be solved in $O(n)$ time.



Weighted Independent Set on Trees

Weighted independent set on trees. Given a tree and node weights $w_v > 0$, find an independent set S that maximizes $\sum_{v \in S} w_v$.




$\text{children}(u) = \{ v, w, x \}$

Weighted Independent Set on Trees

Weighted independent set on trees. Given a tree and node weights $w_v > 0$, find an independent set S that maximizes $\sum_{v \in S} w_v$.

Observation: If (u, v) is an edge such that v is a leaf node, then either OPT includes u , or it includes all leaf nodes incident to u .



$\text{children}(u) = \{ v, w, x \}$

Weighted Independent Set on Trees: DP

Weighted independent set on trees. Given a tree and node weights $w_v > 0$, find an independent set S that maximizes $\sum_{v \in S} w_v$.


Observation: If (u, v) is an edge such that v is a leaf node, then either OPT includes u , or it includes all leaf nodes incident to u .

Dynamic programming solution: Root tree at some node, say r .

- $OPT_{in}(u)$ = max weight independent set rooted at u , containing u .
- $OPT_{out}(u)$ = max weight independent set rooted at u , not containing u .

$$OPT_{in}(u) = w_u + \sum_{v \in \text{children}(u)} OPT_{out}(v)$$

$$OPT_{out}(u) = \sum_{v \in \text{children}(u)} \max \{OPT_{in}(v), OPT_{out}(v)\}$$





$\text{children}(u) = \{v, w, x\}$

Weighted Independent Set on Trees: DP

Theorem: The dynamic programming algorithm find a maximum weighted independent set in trees in $O(n)$ time.

```
Weighted-Independent-Set-In-A-Tree(T) {  
  Root the tree at a node r  
  foreach (node u of T in postorder) {  
    if (u is a leaf) {  
       $M_{in}[u] = w_u$   
       $M_{out}[u] = 0$   
    } else {  
       $M_{in}[u] = \sum_{v \in \text{children}(u)} M_{out}[v] + w_u$   
       $M_{out}[u] = \sum_{v \in \text{children}(u)} \max(M_{out}[v], M_{in}[v])$   
    }  
  }  
  return  $\max(M_{in}[r], M_{out}[r])$   
}
```



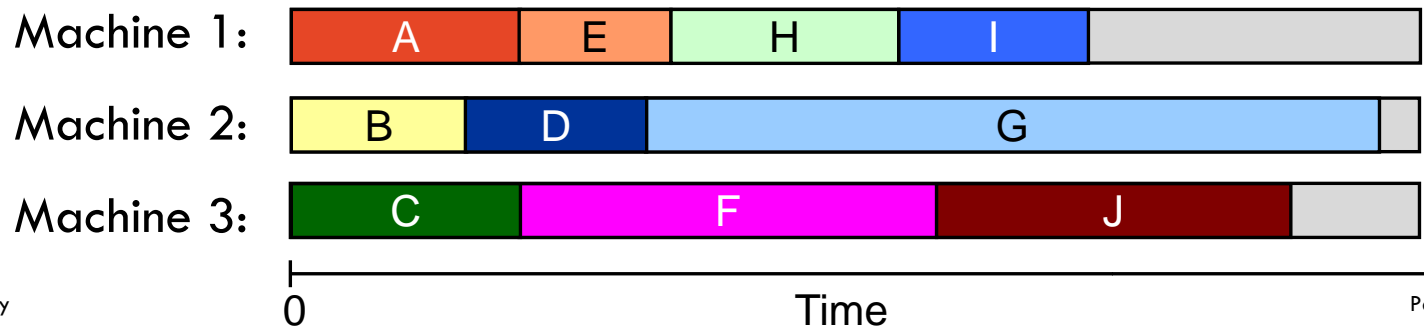
Proof: Takes $O(n)$ time since we visit nodes in postorder and examine each edge exactly once. ▀

11.1 Approximation algorithms: Load Balancing

Load Balancing

Input: m identical machines; n jobs, job j has processing time t_j .

- Job j must run contiguously on one machine.
- A machine can process at most one job at a time.



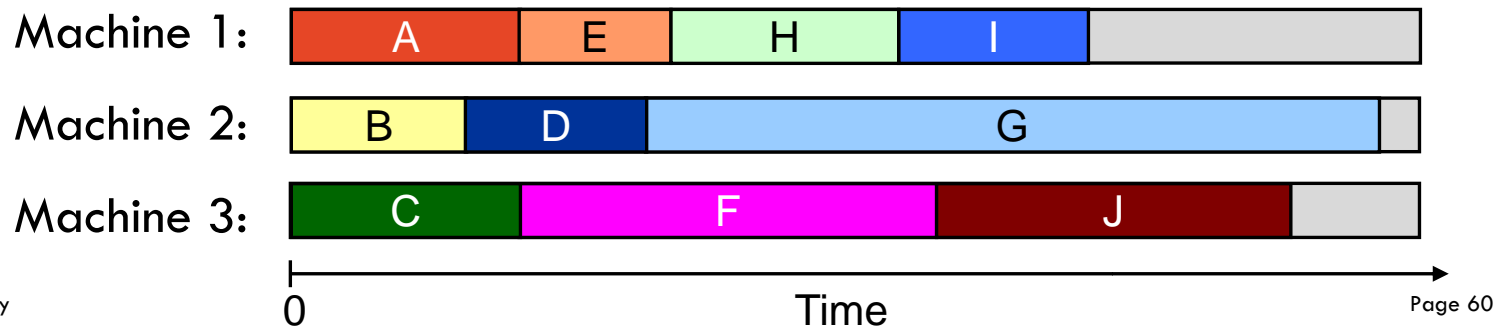
Load Balancing

Input: m identical machines; n jobs, job j has processing time t_j .

- Job j must run contiguously on one machine.
- A machine can process at most one job at a time.

Definition: Let $J(i)$ be the subset of jobs assigned to machine i . The load of machine i is $L_i = \sum_{j \in J(i)} t_j$.

Example: $J(1) = \{A, E, H, I\}$, $J(2) = \{B, D, G\}$, $J(3) = \{C, F, J\}$



Load Balancing

Input: m identical machines; n jobs, job j has processing time t_j .


- Job j must run contiguously on one machine.
- A machine can process at most one job at a time.

Definition: Let $J(i)$ be the subset of jobs assigned to machine i . The load of machine i is $L_i = \sum_{j \in J(i)} t_j$.

Example: $J(1) = \{A, E, H, I\}$, $J(2) = \{B, D, G\}$, $J(3) = \{C, F, J\}$

Definition: The makespan is the maximum load on any machine $L = \max_i L_i$.

Load balancing: Assign each job to a machine to minimize makespan.



Load Balancing: List Scheduling

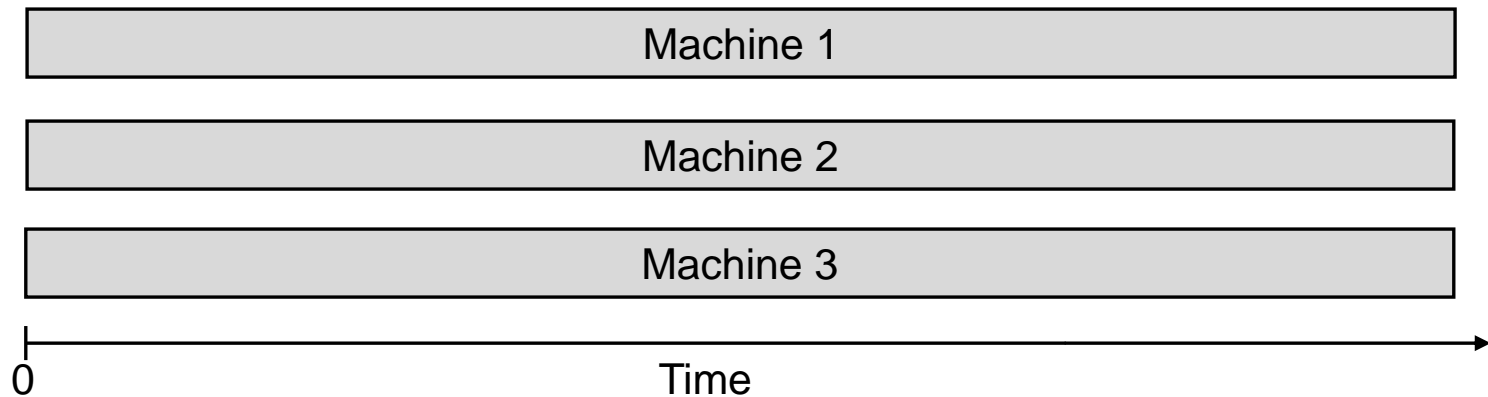

List-scheduling algorithm:

- Consider n jobs in some fixed order.
- Assign job j to machine whose load is smallest so far.

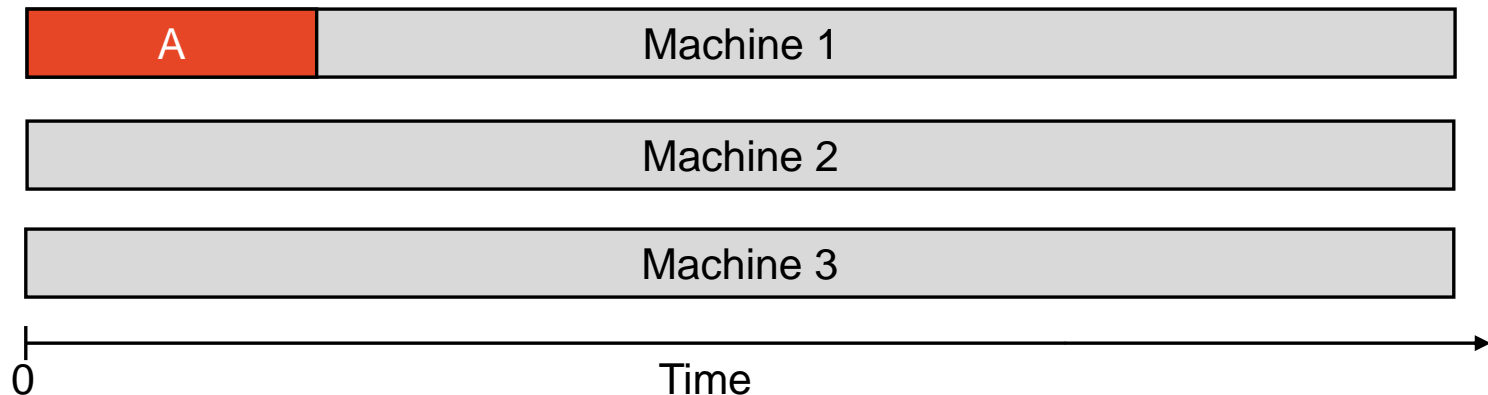

```
List-Scheduling( $m, n, t_1, t_2, \dots, t_n$ ) {  
  for  $i = 1$  to  $m$  {  
     $L_i \leftarrow 0$             $\leftarrow$  load on machine  $i$   
     $J(i) \leftarrow \emptyset$     $\leftarrow$  jobs assigned to machine  $i$   
  }  
  for  $j = 1$  to  $n$  {  
     $i = \operatorname{argmin}_k L_k$        $\leftarrow$  machine  $i$  has smallest load  
     $J(i) \leftarrow J(i) \cup \{j\}$   $\leftarrow$  assign job  $j$  to machine  $i$   
     $L_i \leftarrow L_i + t_j$      $\leftarrow$  update load of machine  $i$   
  }  
}
```

Implementation: $O(n \log n)$ using a priority queue.

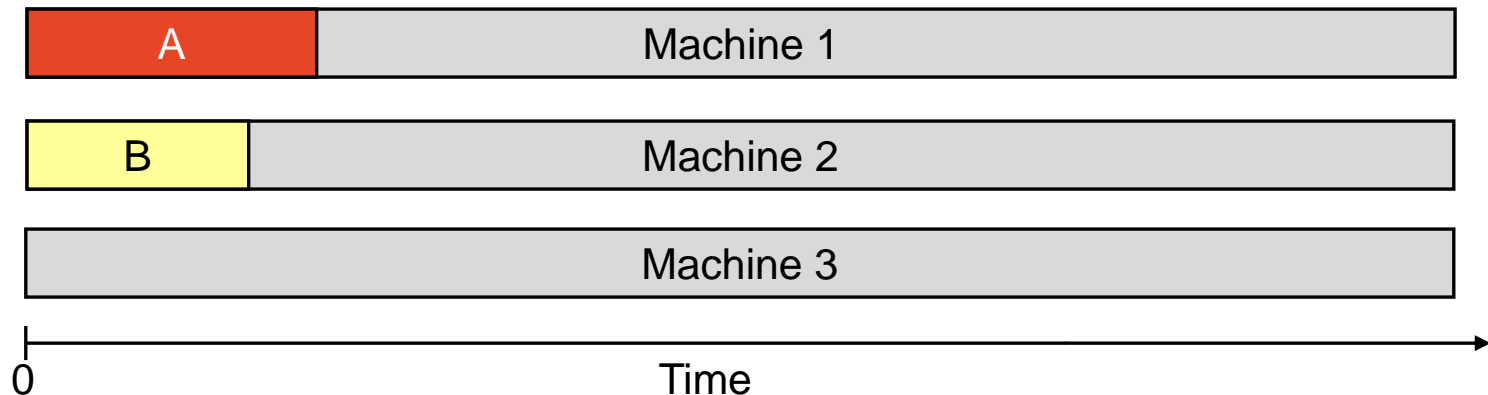

Load Balancing: List Scheduling



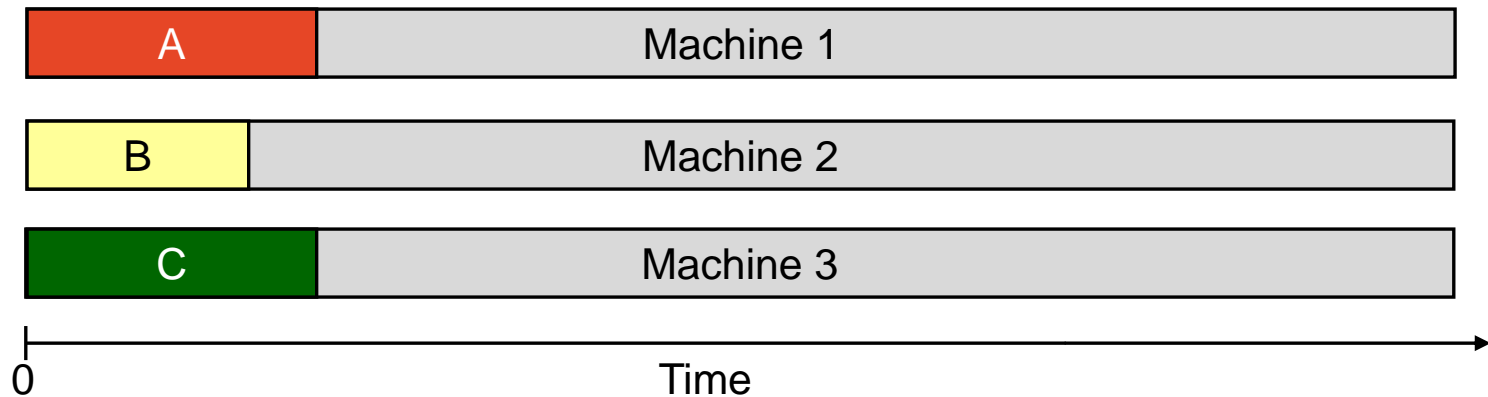

Load Balancing: List Scheduling



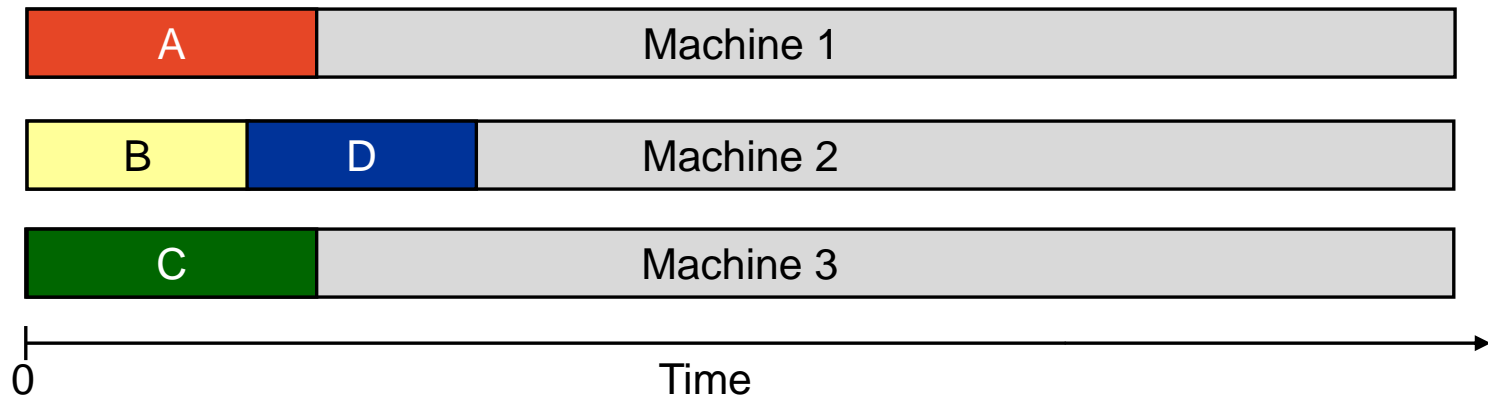

Load Balancing: List Scheduling



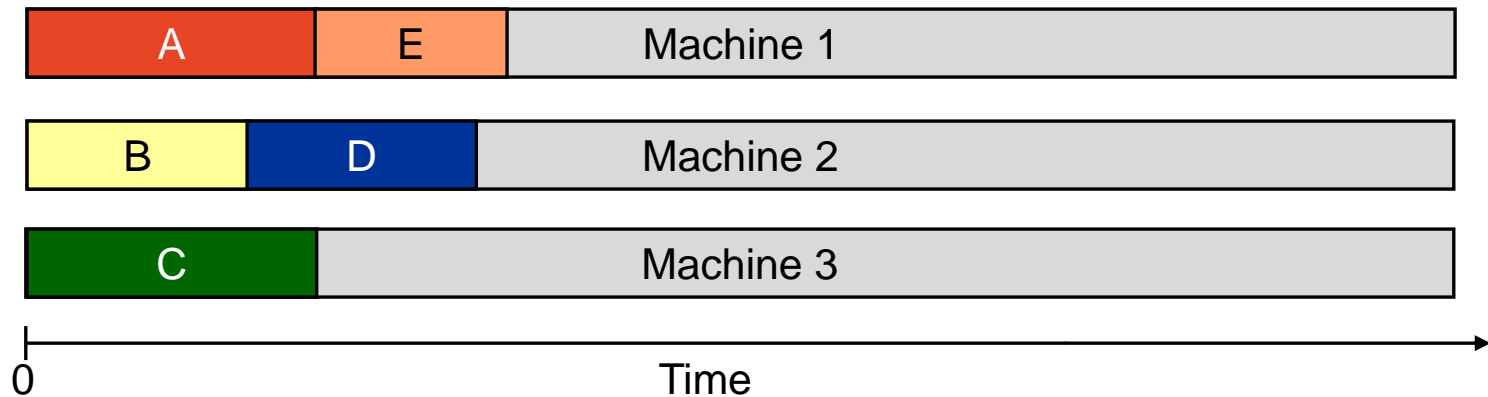

Load Balancing: List Scheduling



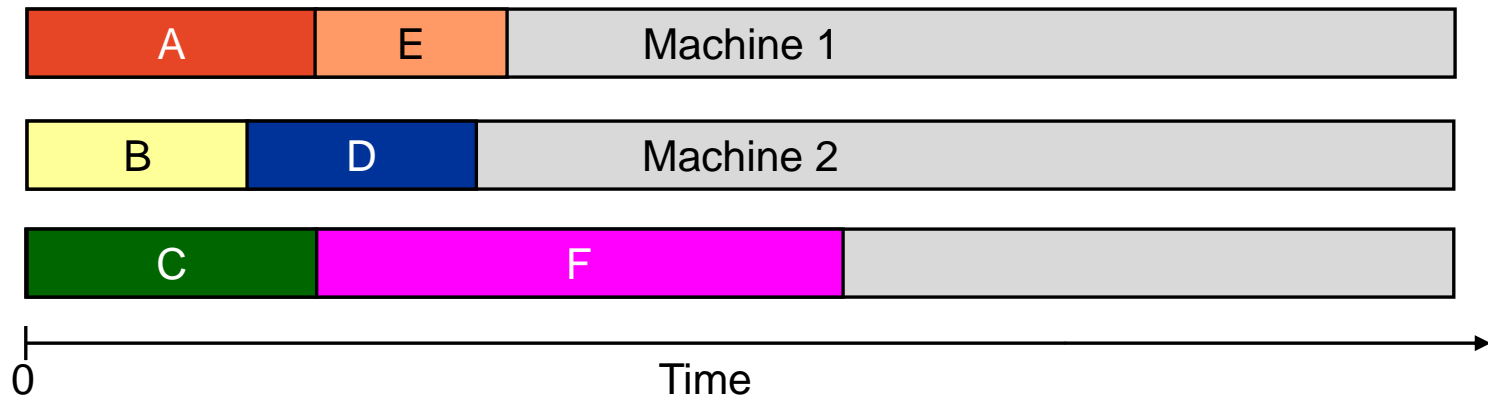
Load Balancing: List Scheduling



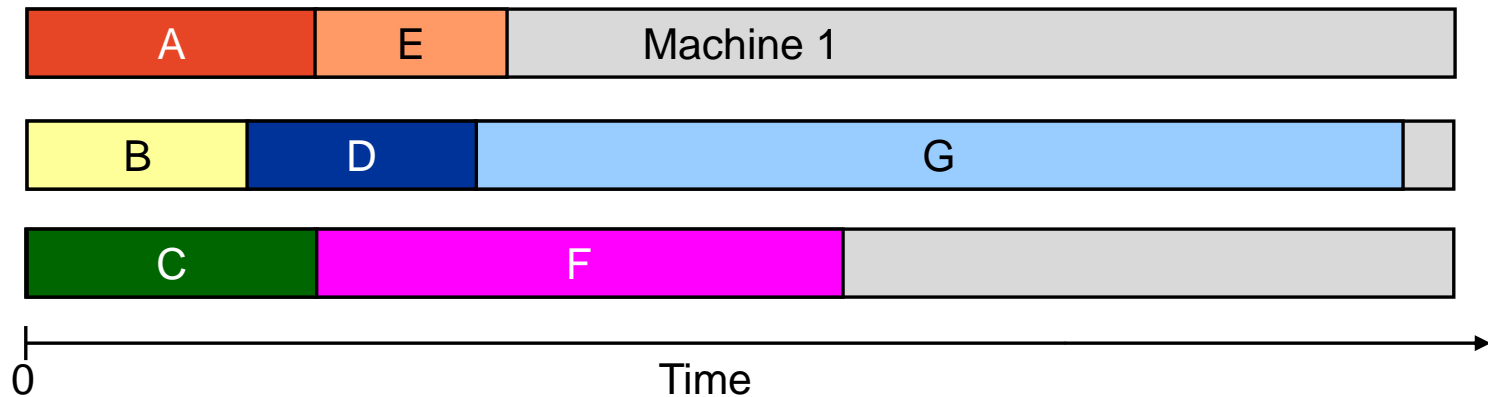
Load Balancing: List Scheduling



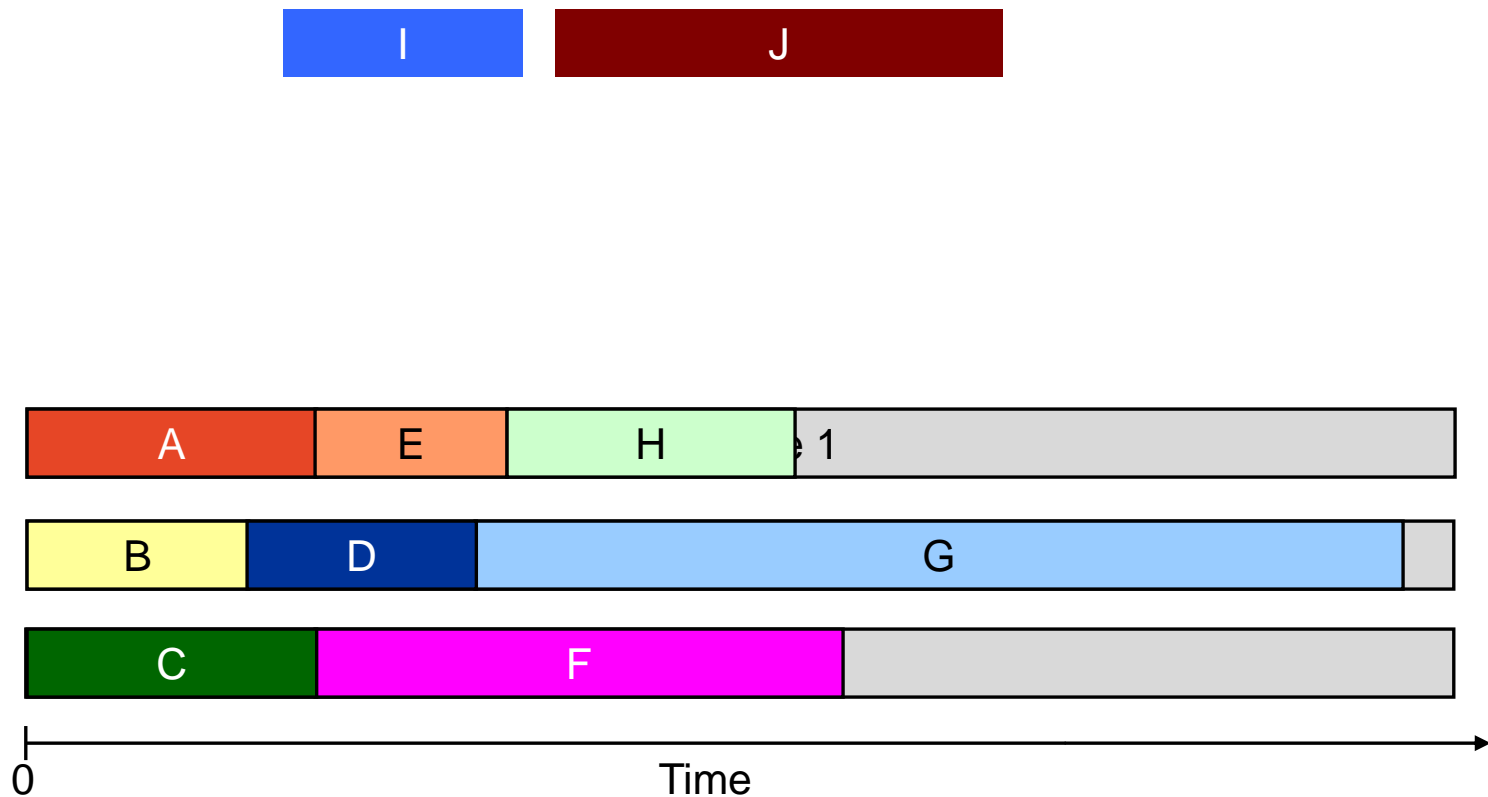
Load Balancing: List Scheduling



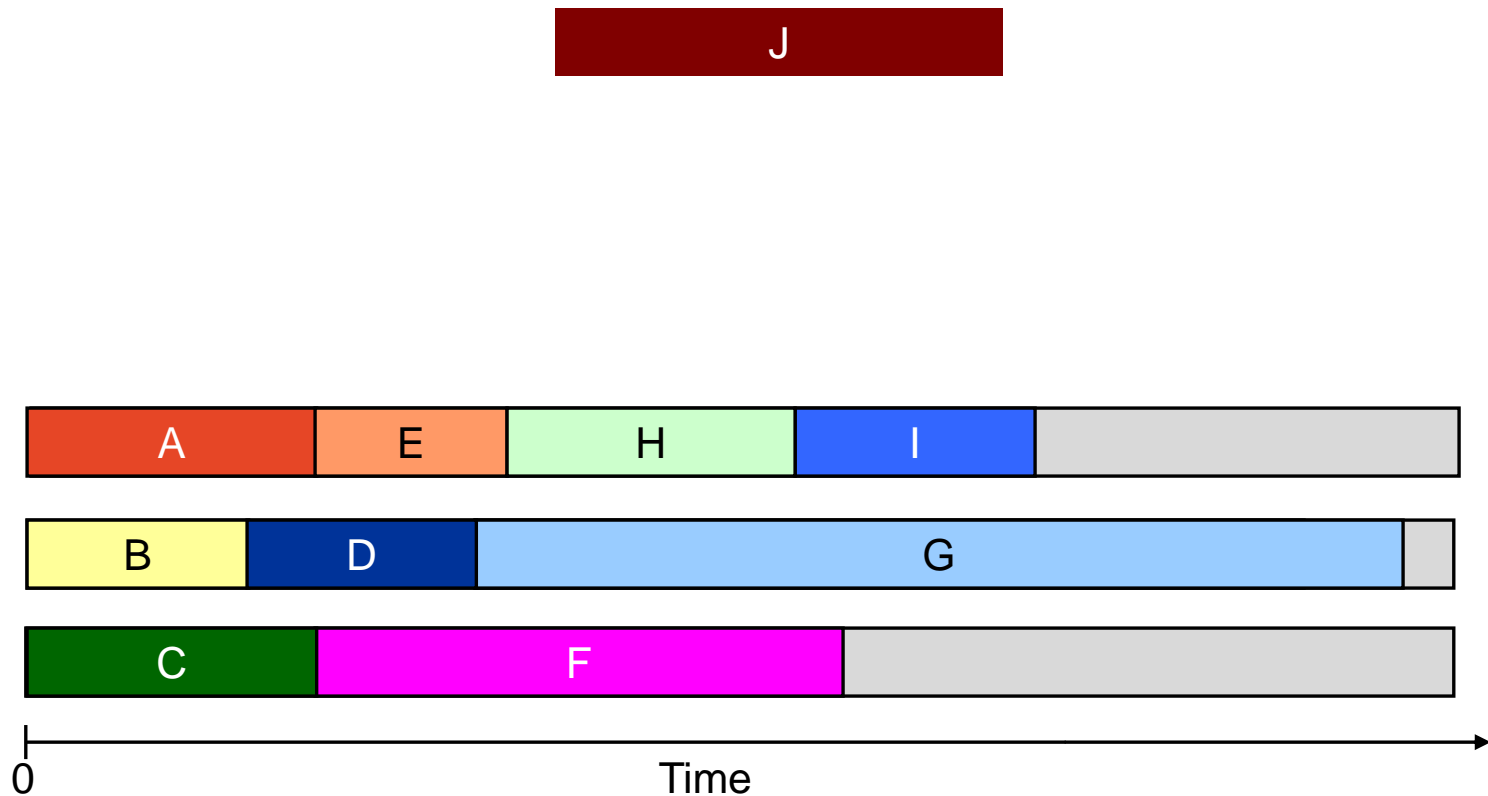
Load Balancing: List Scheduling



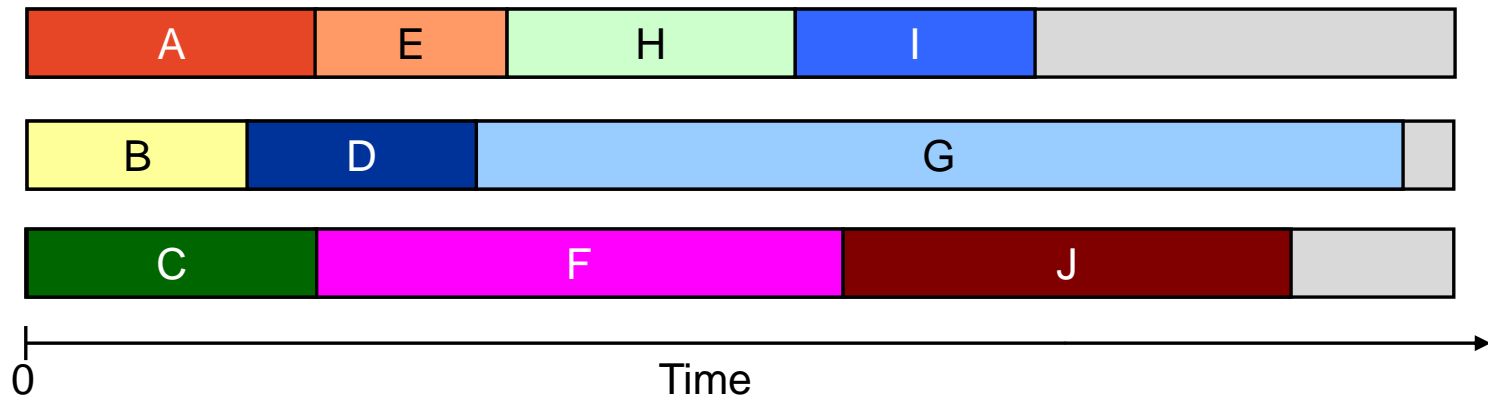
Load Balancing: List Scheduling



Load Balancing: List Scheduling



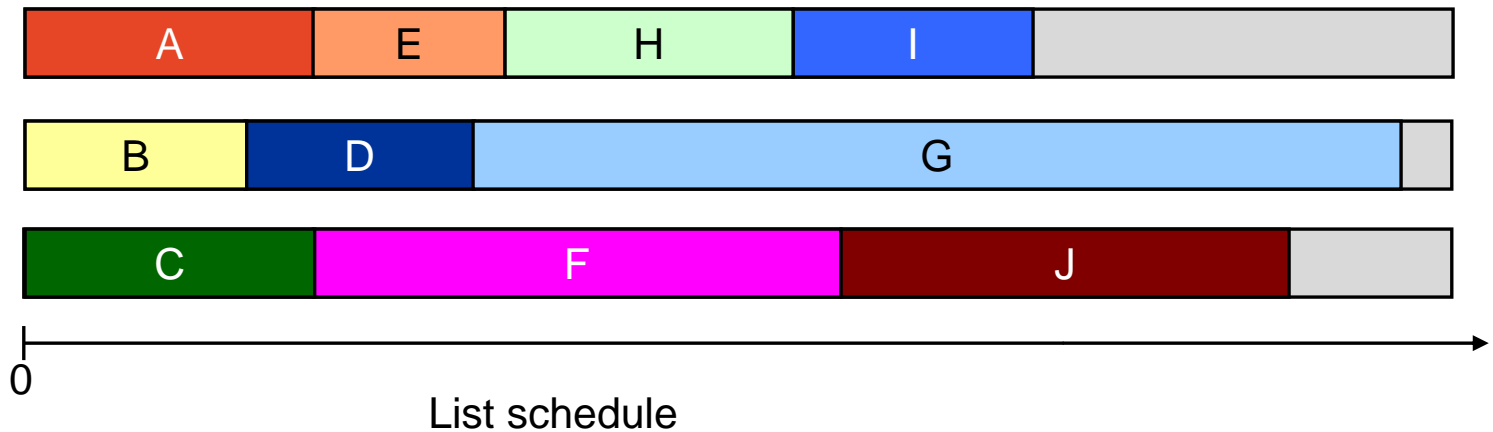
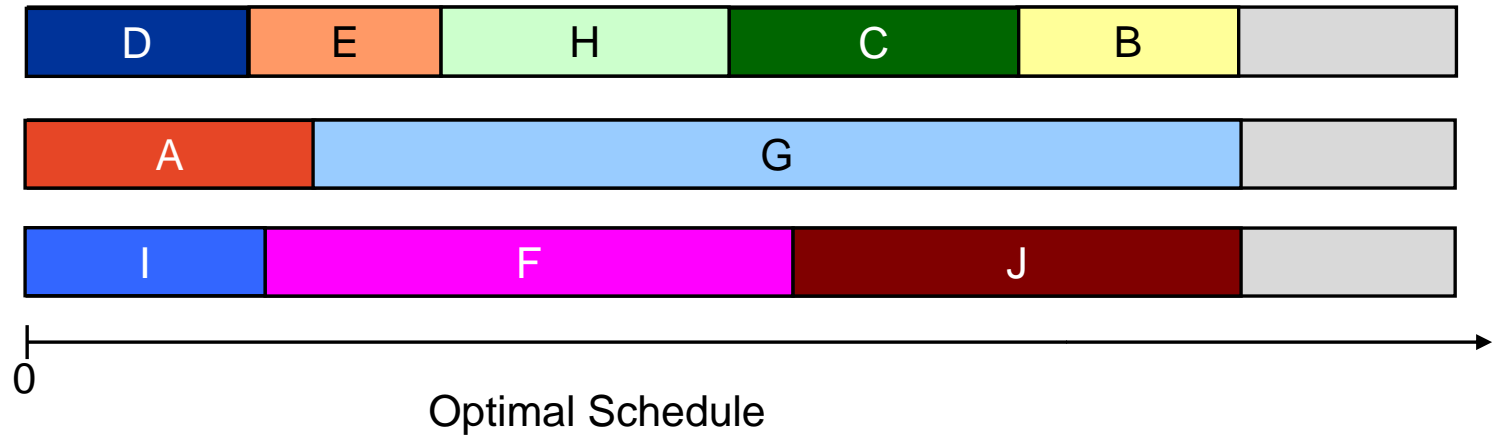
Load Balancing: List Scheduling



Is this a good schedule?

- The schedule may not be optimal (minimum makespan).
- How do we prove that statement?
- We only need to provide a counterexample.

Load Balancing: List Scheduling



How far off can the schedule be from optimal?

Is there an approximation guarantee?

$$\text{Approximation ratio} = \frac{\text{Cost of apx solution}}{\text{Cost of optimal solutions}}$$

An approximation algorithm for a minimization problem requires an approximation guarantee:

- Approximation ratio $\leq c$
- Approximation solution $\leq c \cdot \text{value of optimal solution}$

Load Balancing: List Scheduling Analysis

Theorem: [Graham, 1966]

Greedy algorithm is a 2-approximation.

- First worst-case analysis of an approximation algorithm.
- Need to compare resulting solution with optimal makespan L^* .

Lemma 1: The optimal makespan $L^* \geq \max_i t_i$.

Proof: Some machine must process the most time-consuming job. ▀

Lemma 2: The optimal makespan $L^* \geq \frac{1}{m} \sum_j t_j$.

Proof:


- The total processing time is $\sum_i t_i$.
- One of m machines must do at least a $1/m$ fraction of total work. ▀

Load Balancing: List Scheduling Analysis

Theorem: Greedy algorithm is a 2-approximation.

Proof: Consider load L_i of bottleneck machine i .

- Let j be last job scheduled on machine i .
- When job j assigned to machine i , i had smallest load. Its load before assignment is $L_i - t_j \Rightarrow L_i - t_j \leq L_k$ for all $1 \leq k \leq m$.



Load Balancing: List Scheduling Analysis

Theorem: Greedy algorithm is a 2-approximation.

Proof: Consider load L_i of bottleneck machine i .

- Let j be last job scheduled on machine i .
- When job j was assigned to machine i , i had smallest load. Its load before assignment is $L_i - t_j \Rightarrow L_i - t_j \leq L_k$ for all $1 \leq k \leq m$.
- Sum inequalities over all k and divide by m :

$$\begin{aligned} L_i - t_j &\leq \frac{1}{m} \sum_k L_k \\ &= \frac{1}{m} \sum_k t_k \\ \text{Lemma 1} \rightarrow &\leq L^* \end{aligned}$$

– Now

$$L_i = \underbrace{(L_i - t_j)}_{\leq L^*} + \underbrace{t_j}_{\substack{\leq L^* \\ \uparrow \\ \text{Lemma 2}}} \leq 2L^*.$$

Load Balancing: List Scheduling Analysis

Question: Is our analysis tight?

Answer: Yes...more or less.

Example: m machines, $m(m-1)$ jobs of length 1, one job of length m

$m = 10$

										machine 2 idle
										machine 3 idle
										machine 4 idle
										machine 5 idle
										machine 6 idle
										machine 7 idle
										machine 8 idle
										machine 9 idle
										machine 10 idle


Load Balancing: List Scheduling Analysis

Question: Is our analysis tight?

Answer: Yes...more or less.

Example: m machines, $m(m-1)$ jobs of length 1, one job of length m

$m = 10$



optimal makespan = 10

Summary

NP-complete problems show up in many applications. There are different approaches to cope with it:

- Approximation algorithms
- Restricted cases (trees, bipartite graphs, small solution...)
- Randomized algorithms
- ...

Each approach has its pros and cons.