

# COMP2007/2907 - Algorithms

**Course page:** Blackboard and Piazza (or Ed)

**Lecturer:** Joachim Gudmundsson  
Level 4, Room 416, School of IT  
joachim.gudmundsson@sydney.edu.au  
Ph. 9351 4494

<b>Tutors:</b>	Xavi Holt (TA)	Yilun He
	Natalie Tridgell	Joe Godbehere
	Jessica McBroom	Patrick Eades
	Anton Jurisevic	Shumin Kong
	Gengxing Wang	Hisham Husein
	Mingshen Cai	



### Course book:


J. Kleinberg and E. Tardos  
Algorithm Design  
Addison-Wesley

### Outline:

12 lectures  
5 assignments  
10+1 quizzes  
Exam

### Tutorials:

12 tutorials



- › This unit provides an introduction to the design and analysis of algorithms. Its main aims are
  - (i) learn how to develop algorithmic solutions to computational problem
  - (ii) develop understanding of algorithm efficiency.
  
- › Assumes basic knowledge of discrete math
  - graphs
  - big O notation
  - proof techniques



### Assessment:

- Quizzes 20% (average of best 8 out of 10)
- Each assignment 6% (5 assignments - total 30%)
- Exam 50% (minimum 40% required to pass)

Assignments submitted via Blackboard. Turnitin will be used to check every submission

### Collaboration:

- General ideas - Yes!
- Formulation and writing - No!
- Read [Academic Dishonesty and Plagiarism.](#)

- › There will be **5** homework assignments
- › The objective of these is to teach problem solving skills
- › Each assignment represents **6% of your final mark**. Late submissions will be penalized by 25% of the full marks per day.

For example, say you get 80% on your assignment:

If submitted on time = 4.8

Late but within 24 hours =  $4.8 * 0.75 = 3.6$

Between 24 and 48 hours =  $4.8 * 0.5 = 2.4$

Between 48 and 72 hours =  $4.8 * 0.25 = 1.2$

More than 72 hours =  $4.8 * 0 = 0$



- › The final will be 2.5 hours long. It will consist of 6 problems similar to those seen in the tutorials and assignments
- › The final will test your problem solving skills
- › There is a **40% exam barrier**
- › The final exam represents **50% of your final mark**
- › Our advice is that you work hard on the assignments throughout the semester. It's the best preparation for the final.



- › To get the most out of the tutorial, try to solve as many problems as you can *before* the tutorial. Your tutor is there to help you out if you get stuck, not to lecture.
- › We will post solutions to tutorials (see Ed).

- › **Lecture 1** [Mon 31 July]: Introduction
- › **Lecture 2** [Mon 7 Aug]: Graphs
- › **Lecture 3** [Mon 14 Aug]: Greedy algorithms
- › **Lecture 4** [Mon 21 Aug]: Divide & Conquer algorithms
- › **Lecture 5** [Mon 28 Aug]: Sweepline algorithms
- › **Lecture 6** [Mon 4 Sep]: Dynamic programming: basic techniques
- › **Lecture 7** [Mon 11 Sep]: Dynamic programming: interval scheduling and Bellman-Ford
- › **Lecture 8** [Mon 18 Sep]: Network flows I: Theory

**Mon 25 Sep: University break**

**Mon 2 Oct: Labour Day**

- › **Lecture 9** [Mon 9 Oct]: Network flows II: Applications
- › **Lecture 10** [Mon 16 Oct]: NP and intractability
- › **Lecture 11** [Mon 23 Oct]: Coping with hardness
- › **Lecture 12** [Mon 30 Oct]: Recap



# COMP2007/2907: Algorithms

Algorithms then, and now



- Algorithms can have huge impact
- For example -

A report to the White House from 2010 includes the following.

- Professor Martin Grotschel
    - A benchmark production planning model solved using linear programming would have taken 82 years to solve in 1988, using the computers and the linear programming algorithms of the day.
    - Fifteen years later, in 2003, this same model could be solved in roughly 1 minute, an improvement by a factor of roughly 43 million! [Extreme case, but even the average factor is very high.]
-

- In 2003 there were examples of problems that we can solve 43 million times faster than in 1988
  - This is because of better hardware and better algorithms

- In 1988
    - Intel 386 and 386SX
      - About 275,000 transistors
      - clock speeds of 16MHz, 20MHz, 25MHz, and 33MHz
    - MSDOS 4.0 and windows 2.0
    - VGA
  - In 2003
    - Pentium M
      - About 140 million transistors
      - Up to 2.2 GHz
    - AMD Athlon 64
    - Windows XP
-

- In a report to the White House from 2010 includes the following.
  - Professor Martin Grotschel:
    - A benchmark production planning model solved using linear programming would have taken 82 years to solve in 1988, using the computers and the linear programming algorithms of the day.
    - Fifteen years later, in 2003, this same model could be solved in roughly 1 minute, an improvement by a factor of roughly 43 millions

## Observation:

- Hardware: 1,000 times improvement
  - Algorithms: 43,000 times improvement
-

- Efficient algorithms produce results within available resource limits
  - In practice
    - Low polynomial time algorithms behave well
    - Exponential running times are infeasible except for very small instances, or carefully designed algorithms
  - Performance depends on many obvious factors
    - Hardware
    - Software
    - Algorithm
    - Implementation of the algorithm
  - **This unit:** Algorithms
-

- Efficient algorithms “do the job” the way you want them to...
    - Do you need the exact solution?
    - Are you dealing with some special case and not with a general problem?
    - Is it ok if you miss the right solution sometimes?
-

- Complex, highly sophisticated algorithms can greatly improve performance


But...

- Reasonably good algorithmic solutions that avoid simple, or “lazy” mistakes, can have a much bigger impact!
-



## List of topics

Greedy algorithms  
Divide and conquer  
Sweepline  
Dynamic programming  
Network flows  
Mincut theorems  
Approximation  
Optimization problems




## Introduction: Some Representative Problems

# Four Representative Problems: Interval Scheduling

- › Input. Set of jobs with start times and finish times.
- › Goal. Find **maximum cardinality** subset of mutually compatible jobs.


↑  
jobs don't overlap





# Weighted Interval Scheduling


- › Input. Set of jobs with start times, finish times, and weights.
- › Goal. Find **maximum weight** subset of mutually compatible jobs.





# Bipartite Matching

- › **Input.** Bipartite graph.
- › **Goal.** Find **maximum cardinality** matching.






# Independent Set

- › Input. Graph.
- › Goal. Find **maximum cardinality** independent set.

↑  
subset of nodes such that no two  
joined by an edge




# Four Representative Problems

- › Interval scheduling:  $n \log n$  greedy algorithm.
- › Weighted interval scheduling:  $n \log n$  dynamic programming algorithm.
- › Bipartite matching:  $n^k$  maxflow based algorithm.
- › Independent set: NP-complete.



# Algorithm Analysis & Data Structures





- › **Brute force.** For many non-trivial problems, there is a natural brute force search algorithm that checks every possible solution.
  - Typically takes  $2^N$  time or worse for inputs of size  $N$ .
  - Unacceptable in practice.
- › **Desirable scaling property.** When the input size doubles, the algorithm should only slow down by some constant factor  $C$ .

There exists constants  $c > 0$  and  $d > 0$  such that on every input of size  $N$ , its running time is bounded by  $c N^d$  steps.

- › **Definition:** An algorithm is **poly-time** if the above scaling property holds.

- › **Worst case running time.** Obtain bound on **largest possible** running time of algorithm on input of a given size  $N$ .
  - Generally captures efficiency in practice.
  - Draconian view, but hard to find effective alternative.
  
- › **Average case running time.** Obtain bound on running time of algorithm on **random** input as a function of input size  $N$ .
  - Hard (or impossible) to accurately model real instances by random distributions.
  - Algorithm tuned for a certain distribution may perform poorly on other inputs.

- › **Definition:** An algorithm is **efficient** if its running time is polynomial.
- › **Justification:** It really works in practice!
  - In practice, the poly-time algorithms that people develop almost always have low constants and low exponents.
  - Breaking through the exponential barrier of brute force typically exposes some crucial structure of the problem.
- › **Exceptions.**
  - Some poly-time algorithms do have high constants and/or exponents, and are useless in practice.
  - Some exponential-time (or worse) algorithms are widely used because the worst-case instances seem to be rare.

↑  
simplex method  
Unix grep

**Table 2.1** The running times (rounded up) of different algorithms on inputs of increasing size, for a processor performing a million high-level instructions per second. In cases where the running time exceeds  $10^{25}$  years, we simply record the algorithm as taking a very long time.

	$n$	$n \log_2 n$	$n^2$	$n^3$	$1.5^n$	$2^n$	$n!$
$n = 10$	< 1 sec	< 1 sec	< 1 sec	< 1 sec	< 1 sec	< 1 sec	4 sec
$n = 30$	< 1 sec	< 1 sec	< 1 sec	< 1 sec	< 1 sec	18 min	$10^{25}$ years
$n = 50$	< 1 sec	< 1 sec	< 1 sec	< 1 sec	11 min	36 years	very long
$n = 100$	< 1 sec	< 1 sec	< 1 sec	1 sec	12,892 years	$10^{17}$ years	very long
$n = 1,000$	< 1 sec	< 1 sec	1 sec	18 min	very long	very long	very long
$n = 10,000$	< 1 sec	< 1 sec	2 min	12 days	very long	very long	very long
$n = 100,000$	< 1 sec	2 sec	3 hours	32 years	very long	very long	very long
$n = 1,000,000$	1 sec	20 sec	12 days	31,710 years	very long	very long	very long

- › **Upper bounds.**  $T(n)$  is  $O(f(n))$  if there exist constants  $c > 0$  and  $n_0 \geq 0$  such that for all  $n \geq n_0$  we have  $T(n) \leq c \cdot f(n)$ .
- › **Lower bounds.**  $T(n)$  is  $\Omega(f(n))$  if there exist constants  $c > 0$  and  $n_0 \geq 0$  such that for all  $n \geq n_0$  we have  $T(n) \geq c \cdot f(n)$ .
- › **Tight bounds.**  $T(n)$  is  $\Theta(f(n))$  if  $T(n)$  is both  $O(f(n))$  and  $\Omega(f(n))$ .
- › **Ex:**  $T(n) = 32n^2 + 17n + 32$ .
  - $T(n)$  is  $O(n^2)$ ,  $O(n^3)$ ,  $\Omega(n^2)$ ,  $\Omega(n)$ , and  $\Theta(n^2)$  .
  - $T(n)$  is not  $O(n)$ ,  $\Omega(n^3)$ ,  $\Theta(n)$ , or  $\Theta(n^3)$ .

- › Slight abuse of notation.  $T(n) = O(f(n))$ .
  - Asymmetric:
    - $f(n) = 5n^3$ ;  $g(n) = 3n^3$
    - $f(n) = O(n^3) = g(n)$
    - but  $f(n) \neq g(n)$ .
  - Better notation:  $T(n) \in O(f(n))$ .
  
- › **Meaningless statement.** Any comparison-based sorting algorithm requires at least  $O(n \log n)$  comparisons.
  - Statement doesn't "type-check."
  - Use  $\Omega$  for lower bounds.

### › Transitivity

- If  $f = O(g)$  and  $g = O(h)$  then  $f = O(h)$ .
- If  $f = \Omega(g)$  and  $g = \Omega(h)$  then  $f = \Omega(h)$ .
- If  $f = \Theta(g)$  and  $g = \Theta(h)$  then  $f = \Theta(h)$ .

### › Additivity

- If  $f = O(h)$  and  $g = O(h)$  then  $f + g = O(h)$ .
- If  $f = \Omega(h)$  and  $g = \Omega(h)$  then  $f + g = \Omega(h)$ .
- If  $f = \Theta(h)$  and  $g = O(h)$  then  $f + g = \Theta(h)$ .



# Asymptotic Bounds for Some Common Functions

- › **Polynomials.**  $a_0 + a_1n + \dots + a_dn^d$  is  $\Theta(n^d)$  if  $a_d > 0$ .
- › **Polynomial time.** Running time is  $O(n^d)$  for some constant  $d$  independent of the input size  $n$ .

- › **Logarithms.**  $O(\log_a n) \overset{\uparrow}{=} O(\log_b n)$  for any constants  $a, b > 0$ .  
can avoid specifying the base

- › **Logarithms.** For every  $x > 0$ ,  $\log n \overset{\uparrow}{=} O(n^x)$ .  
log grows slower than every polynomial

- › **Exponentials.** For every  $r > 1$  and every  $d > 0$ ,  $n^d \overset{\uparrow}{=} O(r^n)$ .


every exponential grows faster than every polynomial



- › **Linear time.** Running time is at most a constant factor times the size of the input.
- › Computing the maximum. Compute maximum of  $n$  numbers  $a_1, \dots, a_n$ .

```
max ← a1
for i = 2 to n
{
    if (ai > max)
        max ← ai
}
```

- › **Merge.** Combine two sorted lists  $A = a_1, a_2, \dots, a_n$  with  $B = b_1, b_2, \dots, b_n$  into one sorted list.



```
i = 1, j = 1
while (both lists are nonempty) {
    if (a_i ≤ b_j) then append a_i to output list and
    increment i
    else append b_j to output list and increment j
}
append remainder of nonempty list to output list
```

- ›  $O(n \log n)$  time. Arises in divide-and-conquer algorithms.
- › Sorting. Mergesort and heapsort are sorting algorithms that perform  $O(n \log n)$  comparisons.
- › Largest empty interval. Given  $n$  time-stamps  $x_1, \dots, x_n$  on which copies of a file arrive at a server, what is largest interval of time when no copies of the file arrive?
- ›  $O(n \log n)$  solution. Sort the time-stamps. Scan the sorted list in order, identifying the maximum gap between successive time-stamps.

- › **Quadratic time.** Enumerate all pairs of elements.
- › Closest pair of points. Given a list of  $n$  points in the plane  $(x_1, y_1), \dots, (x_n, y_n)$ , find the pair that is closest.
- ›  $O(n^2)$  solution. Try all pairs of points.

```
min ←  $(x_1 - x_2)^2 + (y_1 - y_2)^2$  ← don't need to  
for i = 1 to n {  
    for j = i+1 to n {  
        d ←  $(x_i - x_j)^2 + (y_i - y_j)^2$   
        if (d < min)  
            min ← d ← see chapter 5  
    }  
}
```

- › **Cubic time.** Enumerate all triples of elements.
- › **Set disjointness.** Given  $n$  sets  $S_1, \dots, S_n$  each of which is a subset of  $1, 2, \dots, n$ , is there some pair of these which are disjoint?
- ›  $O(n^3)$  solution. For each pairs of sets, determine if they are disjoint.

```
foreach set  $S_i$  {  
    foreach other set  $S_j$  {  
        foreach element  $p$  of  $S_i$  {  
            determine whether  $p$  also belongs to  $S_j$   
        }  
        if (no element of  $S_i$  belongs to  $S_j$ )  
            report that  $S_i$  and  $S_j$  are disjoint  
    }  
}
```

# Polynomial Time: $O(n^k)$ Time

- › Independent set of size  $k$ . Given a graph, are there  $k$  nodes such that no two are joined by an edge?  
↖  
k is a constant
- ›  $O(n^k)$  solution. Enumerate all subsets of  $k$  nodes.

```
foreach subset S of k nodes {  
    check whether S is an independent set  
    if (S is an independent set)  
        report S is an independent set  
    }  
}
```

- Check whether  $S$  is an independent set =  $O(k^2)$ .

- Number of  $k$  element subsets =  $\binom{n}{k} = \frac{n(n-1)(n-2)\cdots(n-k+1)}{k(k-1)(k-2)\cdots(2)(1)} \leq \frac{n^k}{k!}$

- $O(k^2 n^k / k!) = O(n^k)$ .

poly-time for  $k=17$ ,  
but not practical

↙

- › Independent set. Given a graph, what is maximum size of an independent set?
- ›  $O(n^2 2^n)$  solution. Enumerate all subsets.

```
S* ← ∅  
foreach subset S of nodes {  
    check whether S is an independent set  
    if (S is largest independent set seen so far)  
        update S* ← S  
}  
}
```

- › You must learn the asymptotic order of growth. It is fundamental when measuring the performance of an algorithm.
  - $O$ -notation
  - $\Omega$ -notation
  - $\Theta$ -notation
- › Transitivity and additivity





# Basic dynamic data structures

Assumed knowledge:


- Linked lists
- Queues
- Stacks
- Balanced binary trees

- › Programs manipulate data
  - › Data should be organized so manipulations will be efficient
    - Search (e.g. Finding a word/file/web page)
  - › Good programs are powered by good data structures
  - › Naïve choices are often much less efficient than clever choices
  - › Data structures are existing tools that can help you
    - guide your design, and
    - save you time (avoid re-inventing the wheel)
-

# The Queue data structure


- › The **Queue** data structure stores arbitrary objects
  - › Insertions and deletions follow the first-in first-out (FIFO) scheme
  - › Insertions are at the rear of the queue and removals are at the front of the queue
  - › Main queue operations:
    - **enqueue**(object): inserts an element at the end of the queue
    - object **dequeue**(): removes and returns the element at the front of the queue
  - › Auxiliary queue operations:
    - object **front**(): returns the element at the front without removing it
    - integer **size**(): returns the number of elements stored
    - boolean **isEmpty**(): indicates whether no elements are stored
-

- › The **Stack** data structure stores arbitrary objects
  - › Insertions and deletions follow the last-in first-out (LIFO) scheme
  - › Think of a spring-loaded plate dispenser
  - › Main stack operations:
    - **push**(object): inserts an element
    - object **pop**(): removes and returns the last inserted element
  - › Auxiliary stack operations:
    - object **top**(): returns the last inserted element without removing it
    - integer **size**(): returns the number of elements stored
    - boolean **isEmpty**(): indicates whether no elements are stored
-



- › Root: node without parent (A)
- › Internal node: node with at least one child (A, B, C, F)
- › External node (a.k.a. leaf ): node without children (E, I, J, K, G, H, D)
- › Ancestors of a node: parent, grandparent, grand-grandparent, etc.
- › Depth of a node: number of ancestors
- › Height of a tree: maximum depth of any node (3)
- › Descendant of a node: child, grandchild, grand-grandchild, etc.

- › Subtree: tree consisting of a node and its descendants






## › Notation

$n$  number of nodes

$e$  number of external nodes


$i$  number of internal nodes

$h$  height



## Properties:

- $e = i + 1$
- $n = 2e - 1$
- $h \leq i$
- $h \leq (n - 1)/2$
- $e \leq 2^h$
- $h \geq \log_2 e$
- $h \geq \log_2 (n + 1) - 1$



- › find is  $O(\log n)$ 
    - height of tree is  $O(\log n)$ , no restructures needed
  
  - › insert is  $O(\log n)$ 
    - initial find is  $O(\log n)$
    - Restructuring up the tree, maintaining heights is  $O(\log n)$
  
  - › remove is  $O(\log n)$ 
    - initial find is  $O(\log n)$
    - Restructuring up the tree, maintaining heights is  $O(\log n)$
-



## › Queues

- Enqueue, dequeue, first and size operations in  $O(1)$  time.

## › Stacks

- Push, pop, top and size operations in  $O(1)$  time

## › Balanced binary trees (e.g. AVL trees)

- Insert, delete and find operations in  $O(\log n)$  time
-