

Algorithms and Complexity

Graphs: Representations and Exploration

Julian Mestre

School of Information Technologies
The University of Sydney



THE UNIVERSITY OF
SYDNEY


Undirected graphs

Let $G=(V,E)$ be an undirected graph:

- V = set of vertices (a.k.a. nodes)
- E = set of edges

Some notation

- $\deg(u)$ = # edges incident on u
- $\deg(G) = \max u \deg(u)$
- $N(u)$ = neighborhood of u
- $\delta(u)$ = edges incident on u
- $n = |V|$
- $m = |E|$




Let $G=(V,E)$ be a directed graph:

- V = set of vertices (a.k.a. nodes)
- E = set of directed edges (a.k.a. arcs)

Some notation

- $\deg^{\text{out}}(u)$ = # arcs out of u
- $\deg^{\text{in}}(u)$ = # arcs into u
- $N^{\text{out}}(u)$ = out neighborhood of u
- $N^{\text{in}}(u)$ = in neighborhood of u




Graphs as a modeling tool

Can model many relations among elements in a set:

- Social network
- Internet topology
- Protein-protein interaction

Can help formulate problems:

- What's the distance between two nodes?
- What's a central node?
- How well connected the network is?
- What's a critical node?



Let $G=(V,E)$ be an undirected graph:

- sequence v_1, v_2, \dots, v_k is a path if (v_i, v_{i+1}) is an edge in E for all $i = 1, \dots, k-1$
- *length* of the path is the # of edges in it
- a path is *simple* if no repeated vertices
- a *cycle* is a path v_1, v_2, \dots, v_k where $v_1 = v_k$
- a *cycle* v_1, v_2, \dots, v_k is simple if v_1, v_2, \dots, v_{k-1} is a simple path
- G is connected if every every vertex can reach every other vertex


We say that G is a *tree* if:

- G is connected and doesn't have a cycle, or equivalently
- G is connected and $|E| = |V| - 1$

Breadth first search

Exploring a graph from a starting vertex **s**, layer by layer:

- $L_0 = \{s\}$
- L_1 = vertices that are one hop away from **s**
- L_2 = vertices that are two hops away from **s**, but not closer
- \vdots
- L_k = vertices that are k hops away from **s**, but not closer



Breadth first search

```
def BFS(G,s):
```

```
    layers = []
```

```
    current_layer = [s]
```

```
    next_layer = []
```

```
    for v in G:
```

```
        seen[v] = false
```

```
    seen[s] = true
```

```
    while "current_layer not empty" :
```

```
        layers.append(current_layer)
```

```
        for u in current_layer:
```

```
            for v in G[u]:
```

```
                if not seen[v]:
```

```
                    next_layer.append(v)
```

```
                    seen[v] = true
```

```
    current_layer = next_layer
```

```
    next_layer = []
```

```
    return layers
```

```
// layer is an empty list
```

```
// current_layer is a list holding s
```

```
// v in G iterates over vertices
```

```
// seen is an associative map
```

```
// G[u] returns the neighbors of u
```

Obs.: Let G be a graph and s be a vertex in G . Suppose $\text{BFS}(G, s)$ returns layers L_0, L_1, \dots, L_k , then:

- if u belongs to some layer L_i , then there is a path from s to u
- if there is a path from s to u , then u belongs to some L_i
- in fact, u belongs to L_i if and only if the shortest s - u path has i edges

Obs.: Edges across layers must connect adjacent layers.

We would like to bound the running time of $\text{BFS}(G, s)$ in terms of the size of G .

We need to decide the implementation details of the data structures used by the algorithm:

- list: `layers`, `current_layer`, `next_layer`
- associative map: `seen`
- graph: `G`

We need an implementation that supports

- iteration over the element in the list
- append at the end

If we use a linked list, we can

- iterate over the elements in $O(n)$ time, where n is the size
- append at the end in $O(1)$ time

If we use an array, we can

- iterate over the elements in $O(n)$ time, where n is the size
- append at the end in $O(1)$ time
- but we end up wasting a lot of space

We need an iteration that supports

- setting/changing key-value pair
- lookup



If we use a Hash table, depending on the implementation

- lookup and setting/changing typically take expected $O(1)$ time
- building or growing the table can take more than linear time

If the keys are $1, 2, \dots, n$, then we can use array

- lookups and setting/changing in $O(1)$ time
- building the table in $O(1)$ time (depends on how memory is allocated)
- uses $\Omega(n)$ space, which is not an issue here

Adjacency lists



Adjacency matrix

	u	v	x	y
u	0	1	0	1
v	1	0	1	1
x	0	1	0	0
y	1	1	0	0

Scan neighborhood of vertex u

- Adj. list : $\Theta(|N(u)|)$
- Adj. matrix : $\Theta(n)$

Check if u and v are adjacent:

- Adj. list : $\Theta(\min\{|N(u)|, |N(v)|\})$
- Adj. matrix : $\Theta(1)$

Space:

- Adj. list : $\Theta(|V|+|E|)$
- Adj. matrix : $\Theta(|V|^2)$

Time complexity of BFS

```
def BFS(G,s):
```

```
    layers = []
```

```
    current_layer = [s]
```

```
    next_layer = []
```

```
    for v in G:
```

```
        seen[v] = false
```

```
    seen[s] = true
```

```
    while "current_layer not empty" :
```

```
        layers.append(current_layer)
```

```
        for u in current_layer:
```

```
            for v in G[u]:
```

```
                if not seen[v]:
```

```
                    next_layer.append(v)
```

```
                    seen[v] = true
```

```
    current_layer = next_layer
```

```
    next_layer = []
```

```
    return layers
```

this loop takes
 $O(|V|)$ time

This loop takes
 $O(|N(u)|)$ time

Adding up over all u , we get
 $O(\sum_u |N(u)|) = O(|E|)$

Graph:



- discrete object encoding a relation between vertices
- representations: adjacency lists, adjacency matrix
- time complexity of basic primitives depends on representation

Breadth first search (BFS):

- a graph exploration strategy
- starting from a vertex s , visit vertices reachable from s , layer by layer
- L_i holds vertices at distance i from s
- Runs in $O(n+m)$ time if the graph is represented with adjacency lists

Depth first search

Pick a starting vertex, follow outgoing edges that lead to new vertices, and backtrack whenever “stuck”.




DFS tree

Depth first search

```
def DFS(G):  
    for u in G  
        visited[u] = false  
        parent[u] = None  
    time = 0  
    for u in G:  
        if not visited[u]:  
            DFS_visit(u)  
    return parent
```

```
def DFS_visit(u):  
    visited[u] = true  
    time = time + 1  
    discovery[u] = time  
    for v in G[u]:  
        if not visited[v]:  
            parent[v] = u  
            DFS_visit(v)  
    time = time + 1  
    finish[u] = time
```



Time complexity of DFS

```
def DFS(G):  
    for u in G  
        visited[u] = false  
        parent[u] = None  
    time = 0  
    for u in G:  
        if not visited[u]:  
            DFS_visit(u)  
    return parent
```

ignoring work done
inside function calls, it
runs in $O(n)$ time

```
def DFS_visit(u):  
    visited[u] = true  
    time = time + 1  
    discovery[u] = time  
    for v in G[u]:  
        if not visited[v]:  
            parent[v] = u  
            DFS_visit(v)  
    time = time + 1  
    finish[u] = time
```

ignoring work done inside
recursive calls, it runs in
 $O(|N(u)|)$ time

$\Rightarrow O(m)$ time
overall here

Obs.: The subset of edges $\{ (u, \text{parent}[u]): u \in V \}$ forms a collection of trees (a.k.a. forest)

Obs.: An undirected graph is connected if and only if we have a single tree in the DFS forest. In fact, each tree corresponds to a connected component of the graph.

Obs.: Each discovery and finish time is a unique number in $[1, 2n]$

Def.: In a connected graph $G=(V,E)$, we say that $(u,v) \in E$ is a cut edge if $(V, E-(u,v))$ is not connected.

Trivial algorithm runs in $O(m^2)$:

- for each edge (u,v) in G , run DFS to check if the new graph is still connected

A better algorithms runs in $O(nm)$:



- run DFS in G
- for each DFS-tree edge (u,v) remove it from G , run DFS to check if the new graph is still connected

Back edges in DFS forest

Def.: If $\text{parent}[u] \neq \text{None}$ then we call $(u, \text{parent}(u))$ a tree edge

Def.: We say that a non-tree edge (u, v) is a back edge if u is a descendant of v in the DFS forest, or vice-versa

Obs.: In the DFS forest every non-tree edge is a back edge



Def.: In a connected graph $G=(V,E)$, we say that $u \in V$ is a cut vertex if $G-u$ is not connected.

Obs.: If u is the root of the DFS tree, then u is a cut vertex if and only if it has two or more children

Obs.: If u is a leaf of the DFS tree, then u is not a cut vertex

Obs.: If u is not the root of the DFS, u is a cut vertex if it has a child v and no vertex in T_v (subtree rooted at v) can “jump over” u

Let u be an internal vertex in the DFS tree.

Def.: $up[u] = \min discovery[v]$ where $v \in N(u)$

Def.: $down\&up[u] = \min up[v]$ where v is in T_u

An internal vertex u is a cut vertex if and only if it has a child v :
 $down\&up[v] = discovery[u]$

Thm.

Given a connect graph, there is an $O(m)$ time algorithm for computing its cut vertices

Another graph exploration strategy that follows edges to new nodes until “stuck”, then backtracks

Vertices are assigned a discovery and a finishing time

In undirected graph, each edge can be a tree edge or a back edge

DFS is useful for solving other graph problems, like cut edges and cut vertices

Runs in $O(n+m)$ time using adjacency lists representation

Quiz I

- during your tutorial session

Tutorial Sheet 2:

- will be posted on Monday 3 August

Assignment I:

- due on Monday 3 August

Assignment 2:

- out on Tuesday 4 August, due on Monday 10 August