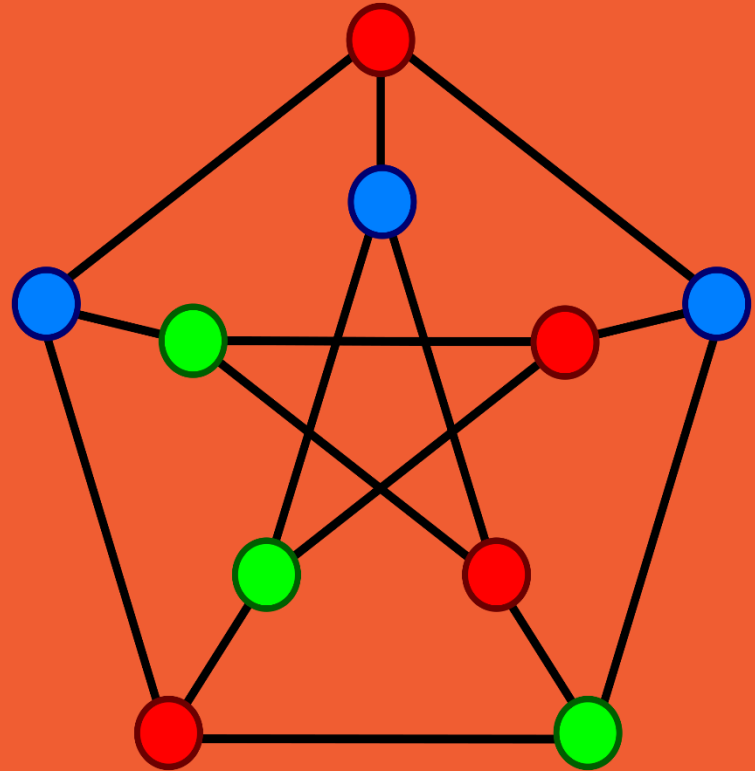


Lecture 2: Graphs (Adv.)

Joachim Gudmundsson




THE UNIVERSITY OF
SYDNEY

3.5 Connectivity in Directed Graphs

Directed Graphs

Directed graph. $G = (V, E)$

- Edge (u, v) goes from node u to node v .



Example. Web graph - hyperlink points from one web page to another.

- Directedness of graph is crucial.
- Modern web search engines exploit hyperlink structure to rank web pages by importance.

Graph Search

Directed reachability. Given a node s , find all nodes reachable from s .

Directed s - t shortest path problem. Given two nodes s and t , what is the length of the shortest path between s and t ?

Graph search. BFS and DFS extend naturally to directed graphs.

Web crawler. Start from web page s . Find all web pages linked from s , either directly or indirectly.

```

def BFS(G,s):

    layers = []
    current_layer = [s]
    next_layer = []
    "mark every vertex except s as not seen"
    while "current_layer not empty" :
        layers.append(current_layer)
        for u in current_layer:
            for v in "neighborhood of u":
                if "haven't seen v yet":
                    next_layer.append(v)
                    "mark v as seen"
            current_layer = next_layer
            next_layer = []

    return layers

```

Strong Connectivity


Definition: Node u and v are **mutually reachable** if there is a path from u to v and also a path from v to u .

Definition: A graph is **strongly connected** if every pair of nodes is mutually reachable.

Lemma: Let s be any node. G is strongly connected iff every node is reachable from s , and s is reachable from every node.

Proof: (\Rightarrow) Follows from definition.

(\Leftarrow) Path from u to v : concatenate u - s path with s - v path.
Path from v to u : concatenate v - s path with s - u path.




Strong Connectivity: Algorithm


Theorem: Can determine if G is strongly connected in $O(m + n)$ time.

Proof:

- Pick any node s .
- Run BFS from s in G .
- Run BFS from s in G^{rev} . ← reverse orientation of every edge in G
- Return true iff all nodes reached in both BFS executions.
- Correctness follows immediately from previous lemma. ▀



strongly connected



not strongly connected

Strong Connectivity

- Consider a graph G and let S_1 and S_2 be two strongly connected components in G of maximal size. Are S_1 and S_2 disjoint?
- Can we compute all the strongly connected components of a graph G efficiently?

Strong Connectivity

Algorithm by Kosaraju 1978 (unpublished)

STRONGLY-CONNECTED-COMPONENTS (G)

1. **Call** DFS(G) to compute finishing times $f[u]$ for all u .
2. **Compute** G^{rev}
3. **Call** DFS(G^{rev}), but in the main loop, consider vertices in order of decreasing $f[u]$ (as computed in first DFS)
4. **Output** the vertices in each tree of the depth-first forest formed in the second DFS as a separate strongly connected component.

Running time: $O(n+m)$

Correctness?


3.6 DAGs and Topological Ordering

Directed Acyclic Graphs (DAGs)


Definition: A **DAG** is a directed graph that contains no directed cycles.

Ex. Precedence constraints: edge (v_i, v_j) means v_i must precede v_j .

Definition: A **topological order** of a directed graph $G = (V, E)$ is an ordering of its nodes as v_1, v_2, \dots, v_n so that for every directed edge (v_i, v_j) we have $i < j$.



a DAG



a topological ordering

Precedence Constraints

Precedence constraints. Edge (v_i, v_j) means task v_i must occur before v_j .

Applications.


- Course prerequisite graph: course v_i must be taken before v_j .
- Compilation: module v_i must be compiled before v_j .
- Pipeline of computing jobs: output of job v_i needed to determine input of job v_j .

Directed Acyclic Graphs

Lemma: If G has a topological order then G is a DAG.

Proof: (by contradiction)

- Suppose that G has a topological order v_1, \dots, v_n and that G also has a directed cycle C . Let's see what happens.
- Let v_i be the lowest-indexed node in C , and let v_j be the node just before v_i in C ; thus (v_j, v_i) is an edge.
- By our choice of i , we have $i < j$.
- On the other hand, since (v_j, v_i) is an edge and v_1, \dots, v_n is a topological order, we must have $j < i$, a contradiction. ■



the supposed topological order: v_1, \dots, v_n

Directed Acyclic Graphs

Lemma: If G has a topological order then G is a DAG.

Question: Does every DAG have a topological ordering?


Question: If so, how do we compute one?

Directed Acyclic Graphs


Lemma: If G is a DAG then G has a node with no incoming edges.

Proof: (by contradiction)

- Suppose that G is a DAG and every node has at least one incoming edge. Let's see what happens.
- Pick any node v , and begin following edges backward from v . Since v has at least one incoming edge (u, v) we can walk backward to u .
- Then, since u has at least one incoming edge (x, u) , we can walk backward to x .
- Repeat until we visit a node, say w , twice.
- Let C denote the sequence of nodes encountered between successive visits to w . C is a cycle. ▀




Topological Ordering Algorithm: Example




Topological order:

Topological Ordering Algorithm: Example




Topological order: v_1

Topological Ordering Algorithm: Example




Topological order: v_1, v_2

Topological Ordering Algorithm: Example




Topological order: v_1, v_2, v_3

Topological Ordering Algorithm: Example



Topological order: v_1, v_2, v_3, v_4

Topological Ordering Algorithm: Example





Topological order: v_1, v_2, v_3, v_4, v_5

Topological Ordering Algorithm: Example



Topological order: $v_1, v_2, v_3, v_4, v_5, v_6$

Topological Ordering Algorithm: Example



Topological order: $v_1, v_2, v_3, v_4, v_5, v_6, v_7$.

Directed Acyclic Graphs

Lemma: If G is a DAG then G has a topological ordering.

Proof: (by induction on n)


- Base case: true if $n = 1$.
- Given DAG on $n > 1$ nodes, find a node v with no incoming edges.
- $G - \{v\}$ is a DAG, since deleting v cannot create cycles.
- By inductive hypothesis, $G - \{v\}$ has a topological ordering.
- Place v first in topological ordering; then append nodes of $G - \{v\}$ in topological order. This is valid since v has no incoming edges. ▀

To compute a topological ordering of G :

Find a node v with no incoming edges and order it first

Delete v from G

Recursively compute a topological ordering of $G - \{v\}$
and append this order after v




Topological Sorting Algorithm: Running Time

Theorem: Algorithm finds a topological order in $O(m + n)$ time.

Proof:

- Maintain the following information:
 - $\text{count}[w] =$ remaining number of incoming edges
 - $S =$ set of remaining nodes with no incoming edges
- Initialization: $O(m + n)$ via single scan through graph.
- Update: to delete v
 - remove v from S
 - decrement $\text{count}[w]$ for all edges from v to w , and add w to S if $\text{count}[w]$ hits 0
 - this is $O(1)$ per edge



Summary: Graphs

- Connectivity in directed graphs
- DAGs
- Topological sort